

COMP9313: Big Data Management



Lecturer: Xubo Wang

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 5.2: Spark IV

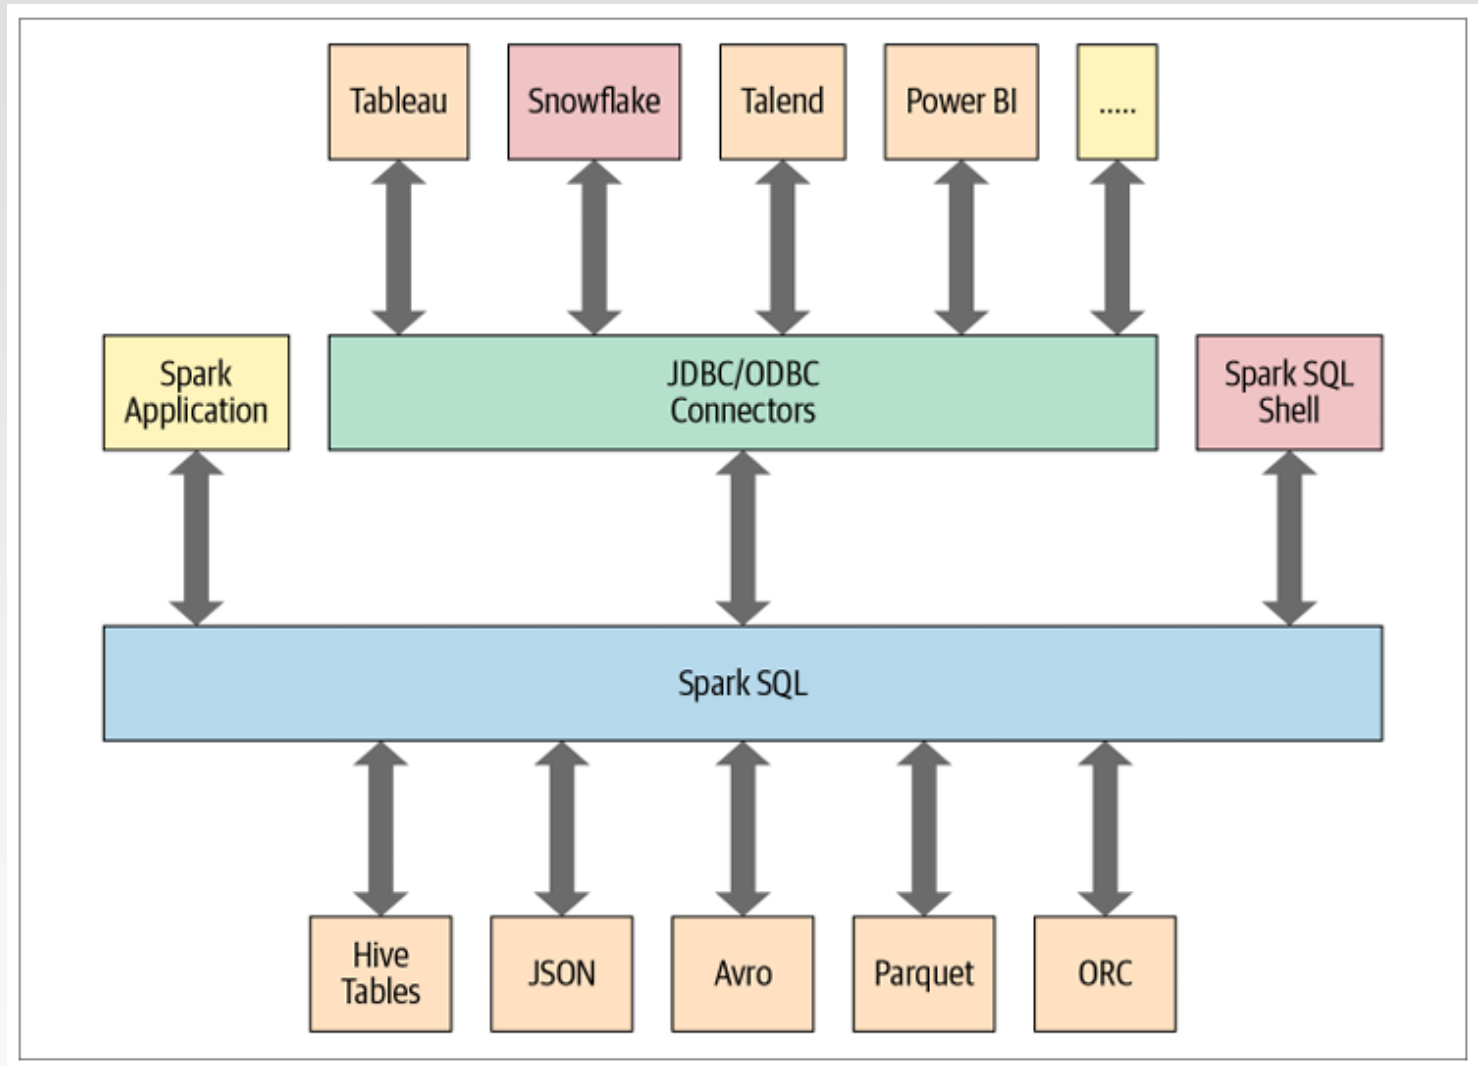


Part 1: Spark SQL

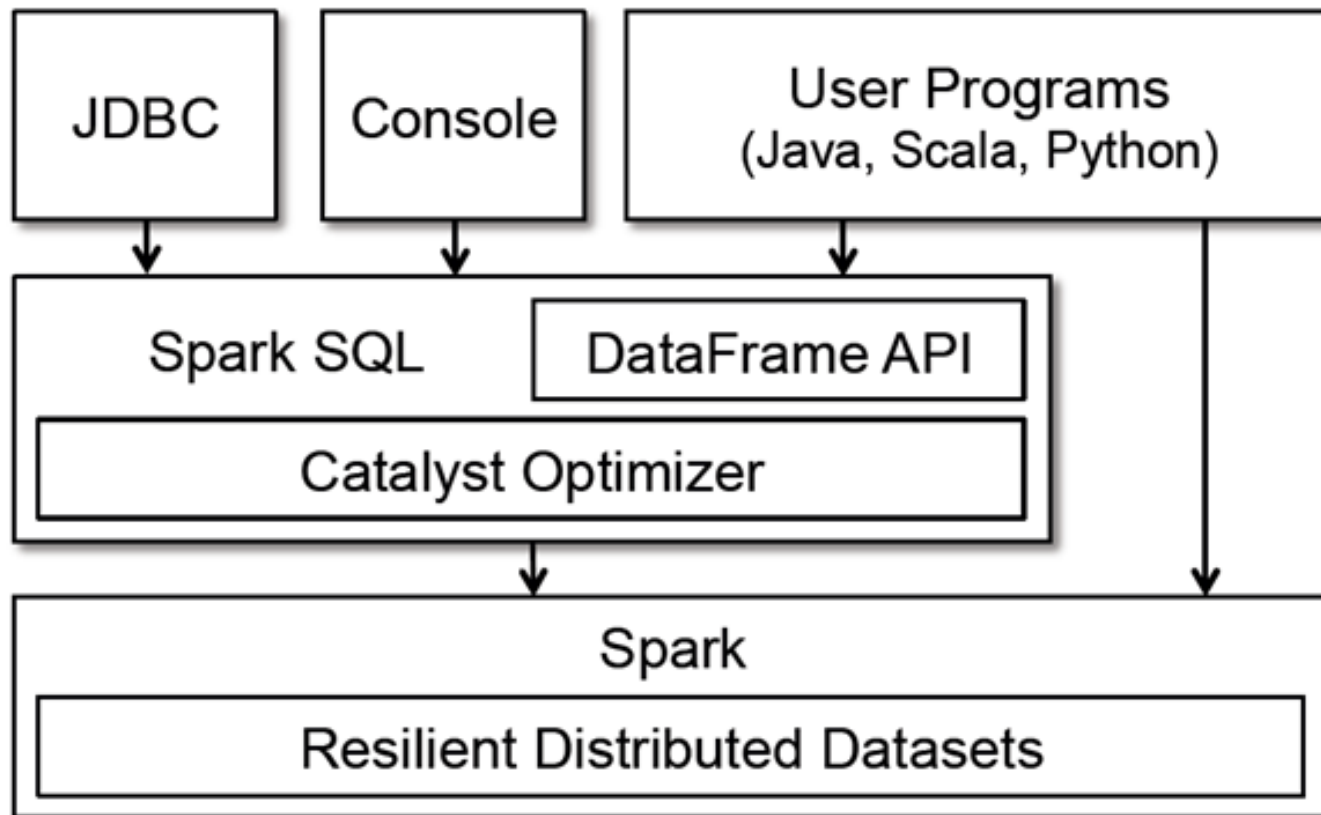
Spark SQL Overview

- ❖ Part of the core distribution since Spark 1.0, Transform RDDs using SQL in early versions (April 2014)
- ❖ Tightly integrated way to work with structured data (tables with rows/columns)
- ❖ Data source integration: Hive, Parquet, JSON, and more
- ❖ Spark SQL is not about SQL.
 - Aims to Create and Run Spark Programs Faster

Spark SQL connectors and data sources



Spark Programming Interface



Starting Point: SparkSession

- ❖ The entry point into all functionality in Spark is the SparkSession class

- Python

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.master("local").appName("Spark SQL
basic example").getOrCreate()
```

- SparkSession since Spark 2.0 provides built-in support for Hive features including the ability to write queries using HiveQL, access to Hive UDFs, and the ability to read data from Hive tables

Creating DataFrames from JSON

- ❖ With a SparkSession, applications can create DataFrames based on the content of a JSON file:

```
df = spark.read.json("examples/src/main/resources/people.json")

// Displays the content of the DataFrame to stdout

df.show()
// +-----+-----+
// |  age|   name|
// +-----+-----+
// |null|Michael|
// |  30|   Andy|
// |  19|  Justin|
// +-----+-----+
```


Running SQL Queries Programmatically

- ❖ The `sql` function on a `SparkSession` enables applications to run SQL queries programmatically and returns the result as a `DataFrame`.

```
// Register the DataFrame as a SQL temporary view
df.createOrReplaceTempView("people")

sqlDF = spark.sql("SELECT * FROM people")
sqlDF.show()
// +----+-----+
// | age|   name|
// +----+-----+
// |null|Michael|
// |  30|   Andy|
// |  19|  Justin|
// +----+-----+
```

Global Temporary View

- ❖ Temporary views in Spark SQL are session-scoped and will disappear if the *session* that creates it terminates
- ❖ Global temporary view: a temporary view that is shared among all sessions and keep alive until the Spark *application* terminates
- ❖ Global temporary view is tied to a system preserved database `global_temp`, and we must use the qualified name to refer it, e.g.
`SELECT * FROM global_temp.view1`

Global Temporary View Example

```
// Register the DataFrame as a global temporary view
df.createGlobalTempView("people")

// Global temporary view is tied to a system preserved database `global_temp`
spark.sql("SELECT * FROM global_temp.people").show()
// +----+-----+
// | age|    name|
// +----+-----+
// |null|Michael|
// |  30|    Andy|
// |  19|   Justin|
// +----+-----+

// Global temporary view is cross-session
spark.newSession().sql("SELECT * FROM global_temp.people").show()
// +----+-----+
// | age|    name|
// +----+-----+
// |null|Michael|
// |  30|    Andy|
// |  19|   Justin|
// +----+-----+
```

Find full example code at

<https://github.com/apache/spark/blob/master/examples/src/main/scala/org/apache/spark/examples/sql/SparkSQLExample.scala>

Spark SQL Built-in Functions

- ❖ Spark SQL provides several built-in standard functions `org.apache.spark.sql.functions` to work with `DataFrame/Dataset` and SQL queries. All these Spark SQL Functions return `org.apache.spark.sql.Column` type.
 - String Functions
 - Date & Time Functions
 - Collection Functions
 - Math Functions
 - Aggregate Functions
 - Window Functions
 - You can check the examples of these functions at:
<https://spark.apache.org/docs/latest/api/python/reference/pyspark.sql/functions.html>
 - In order to use these SQL Standard Functions, you need to import below packing into your application.

```
from pyspark.sql.functions import *
```

WordCount using Spark SQL

❖ Create temporary view

```
fileDF = spark.read.text("file:///home/comp9313/inputText")

from pyspark.sql.functions import *
wordsDF = fileDF.select(explode(split(fileDF.value, ' ')).alias("word"))
wordsDF.createOrReplaceTempView("wc")
```

❖ Count words using SQL

```
wc_df = spark.sql("SELECT word, COUNT(*) as count FROM wc GROUP BY  
word ORDER BY count DESC")
wc_df.show()
```

PySpark Standalone Code (RDD)

```
from pyspark import SparkContext, SparkConf
import sys

class wordCount:
    def run(self, inputPath, outputPath):
        conf = SparkConf().setAppName("word count").setMaster("local[3]")
        sc = SparkContext(conf=conf)

        fileRDD = sc.textFile(inputPath)
        wordsRDD = fileRDD.flatMap(lambda line: line.lower().split())
        pairsRDD = wordsRDD.map(lambda word: (word, 1))
        countRDD = pairsRDD.reduceByKey(lambda a, b: a+b)

        countRDD.saveAsTextFile(outputPath)
        sc.stop()

if __name__ == "__main__":
    if len(sys.argv) != 3:
        print("wrong inputs")
        sys.exit(-1)
    wordCount().run(sys.argv[1], sys.argv[2])
```

PySpark Standalone Code (RDD)

- ❖ The input and output could be on HDFS or on your local file system
 - We receive them from the command line
 - ▶ `spark-submit wordcount.py file:///home/comp9313/pg100.txt file:///home/comp9313/outuput`
 - ▶ `spark-submit wordcount.py hdfs:///localhost:9000/user/comp9313/input hdfs:///localhost:9000/user/comp9313/output`
- ❖ We can use the RDD `textFile()` operation to read the data into an RDD
- ❖ We can use RDD `saveAsTextFile()` operation to write the data to disk
 - The result contains parentheses by default

```
('school', 51)
('20210119,nt', 1)
('law', 19)
('darwin', 22)
('20160326,body', 1)
('believed', 4)
('nurse', 9)
```

- You can format the output and then save to file
- ❖ Remember to release the resources by `sc.stop()` finally

PySpark Standalone Code (DataFrame)

```
from pyspark.sql.session import SparkSession
from pyspark.sql.functions import *
import sys

class wordCount:
    def run(self, inputPath, outputPath):
        spark = SparkSession.builder.master("local").appName("word
count").getOrCreate()
        fileDF = spark.read.text(inputPath)
        wordsDF = fileDF.selectExpr("explode(split(value, ' ')) as
word").withColumn("word", lower(col("word")))

        countDF = wordsDF.groupBy("word").count()
        countDF.write.format("csv").save(outputPath)
        #resDF = countDF.select(concat(col("word"), lit(","), col("count")))
        #resDF.write.text(outputPath)
        spark.stop()

if __name__ == "__main__":
    if len(sys.argv) != 3:
        print("wrong inputs")
        sys.exit(-1)
    wordCount().run(sys.argv[1], sys.argv[2])
```


PySpark Standalone Code (DataFrame)

- ❖ We first need to create a SparkSession object spark
 - You can utilize the SparkSession.builder to configure your task:
`spark = SparkSession.builder.master("local").appName("word count").getOrCreate()`
- ❖ The input and output could be on HDFS or on your local file system
- ❖ Remember to release the resources by `spark.stop()` finally

PySpark Standalone Code (DataFrame)

- ❖ We can use the `pyspark.sql.DataFrameReader.text()` operation to read the text data into a DataFrame
 - `pyspark.sql.DataFrameReader.csv()`
 - `pyspark.sql.DataFrameReader.json()`
 - <https://spark.apache.org/docs/latest/api/python/reference/pyspark.sql/io.html>
- ❖ We can use `pyspark.sql.DataFrameWriter.text()` to write a DataFrame **with a single column of string type** to a file
- ❖ We can use `pyspark.sql.DataFrameWriter.csv()` or `pyspark.sql.DataFrameWriter.format("csv").save()` to store the data as a csv file
 - You can also use other formats such as json
 - You can also use `pyspark.sql.DataFrameWriter.format("text").save()`, but it also requires a DataFrame **with a single column of string type**

select() vs selectExpr()

- ❖ Both functions are used to select columns from a DataFrame.
 - select(): You can select the single or multiple columns of the DataFrame by passing the column names to the select() function

```
df.select("firstname","lastname").show()  
df.select(df.firstname,df.lastname).show()  
df.select(df["firstname"],df["lastname"]).show()
```

#By using col() function

```
from pyspark.sql.functions import col  
df.select(col("firstname"),col("lastname")).show()
```

#By using column index

```
df.select(df.columns[2:4]).show(3)
```

- ❖ selectExpr() projects a set of SQL expressions and returns a new DataFrame. The difference is that it takes a set of SQL expressions in a string to execute.

select() vs selectExpr()

- ❖ Both functions are used to select columns from a DataFrame.
 - select(): You can select the single or multiple columns of the DataFrame by passing the column names to the select() function

```
df.select("firstname", "lastname").show()
df.select(df.firstname, df.lastname).show()
df.select(df["firstname"], df["lastname"]).show()
```

#By using col() function

```
from pyspark.sql.functions import col
df.select(col("firstname"), col("lastname")).show()
```

#By using column index

```
df.select(df.columns[2:4]).show(3)
```

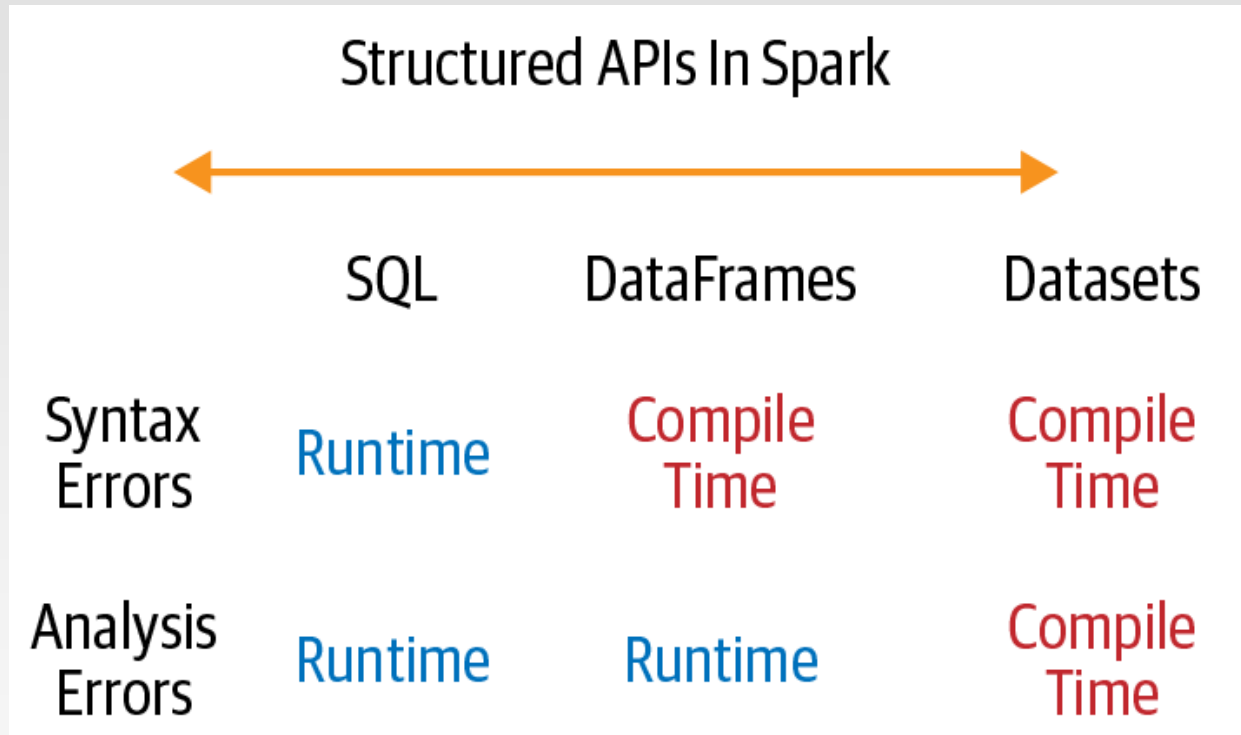
- ❖

```
>>> df = spark.createDataFrame([
    (2, "Alice"), (5, "Bob")], schema=["age", "name"])
>>> df.selectExpr("age * 2", "abs(age)").show()
```

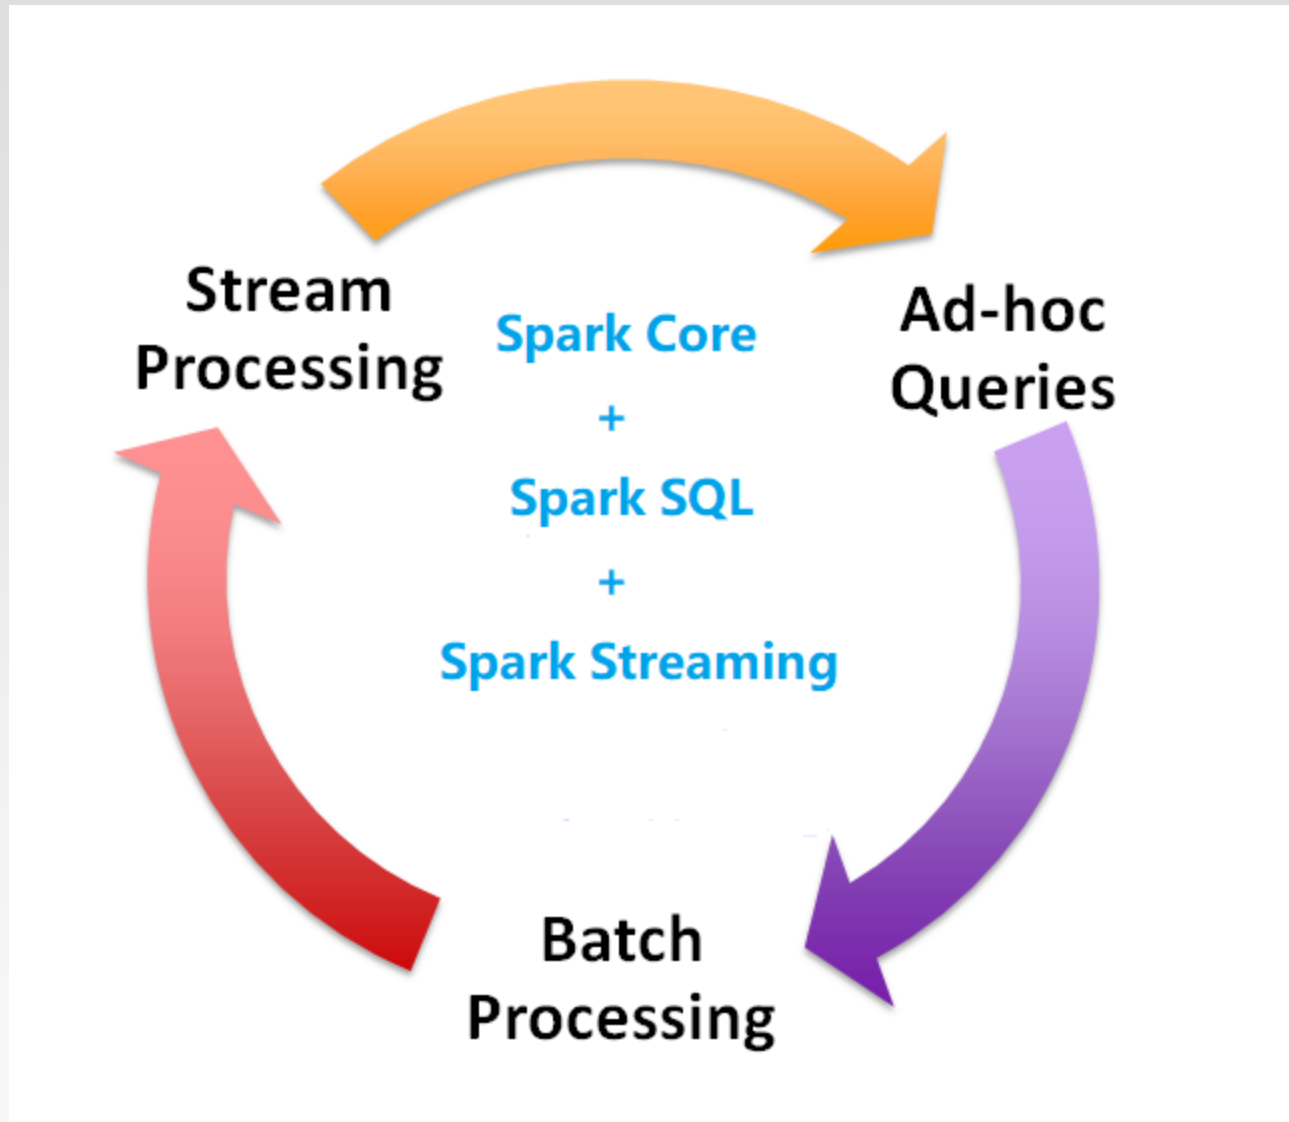
```
+-----+-----+
| (age * 2) | abs(age) |
+-----+-----+
|         4 |         2 |
|        10 |         5 |
+-----+-----+
```

Error Detection of Structured APIs

- ❖ If you want errors caught during compilation rather than at runtime, choose the appropriate API



Vision - one stack to rule them all



Part 2: Spark Programming Practice

Practice

- ❖ Problem 1: Given the data in format of key-value pairs <Int, Int>, find the maximum value for each key across all values associated with that key.

```
pairs = sc.parallelize([(1, 2), (3, 4),... ...])
```

```
resMax = pairs.reduceByKey(lambda x,y: x if x>y else y)
```

```
resMax.foreach(lambda x: print(x[0], x[1]))
```


Practice

- ❖ Problem 2: Given a pair RDD of type [(String, Int)], compute the per-key average

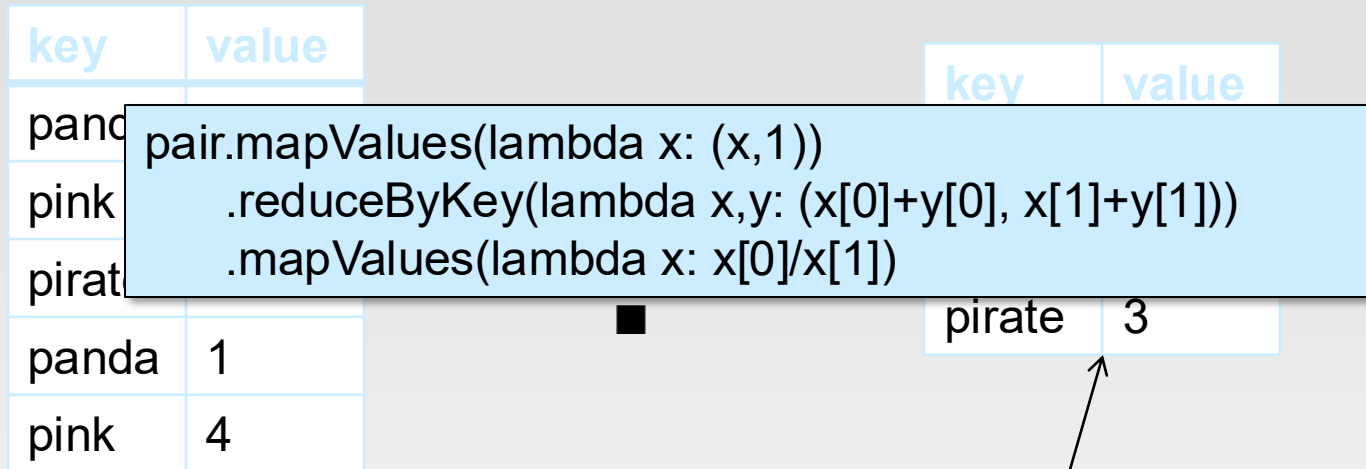
key	value
panda	0
pink	3
pirate	3
panda	1
pink	4



key	value
panda	0.5
pink	3.5
pirate	3

Practice

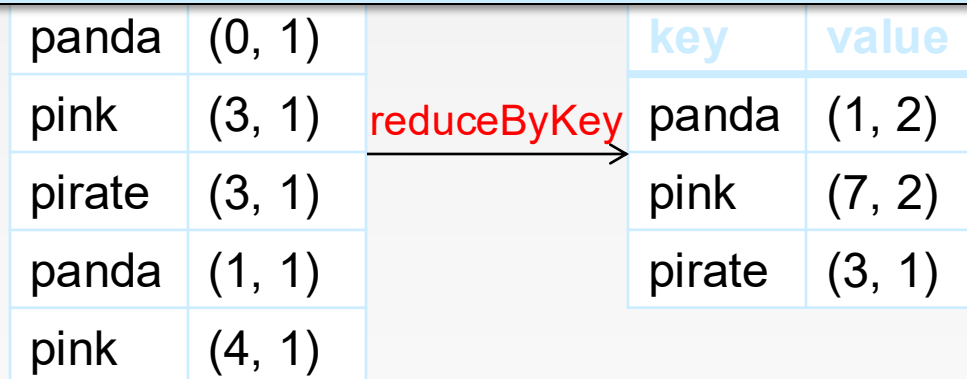
- ❖ Problem 2: Given a pair RDD of type [(String, Int)], compute the per-key average



mapValues →

```
df = spark.createDataFrame(list, schema = "key string, value int")
df.groupBy("key").avg("value").show()
```

mapValues →



Practice (RDD)

- ❖ Problem 3: Given a collection of documents, compute the average length of words starting with each letter.

```
textFile = sc.textFile(inputFile)
words = textFile.flatMap(lambda line: line.split(" ")).map(lambda x: x.lower())

counts = words.filter(lambda x: len(x) >= 1 and x[0] <= 'z' and x[0] >= 'a').map(lambda x:
(x[0], (len(x), 1)))

avgLen = counts.reduceByKey(lambda a, b: (a[0]+b[0], a[1]+b[1])).map(lambda x:
(x[0], x[1][0]/x[1][1]))

avgLen.foreach(lambda x: print(x[0], x[1]))
```

Practice (DataFrame)

- ❖ Problem 3: Given a collection of documents, compute the average length of words starting with each letter.

```
fileDF = spark.read.text(inputFile)

wordsDF = fileDF.select(explode(split(fileDF.value, " ")).alias("word")).withColumn("word",
lower(col("word")))

wordsDF = wordsDF.filter(length(col("word")) >= 1).filter((col("word").substr(0,1)<= 'z') &
(col("word").substr(0,1)>='a'))

pairDF = wordsDF.select(wordsDF.word.substr(0, 1),length(wordsDF.word)).toDF("letter",
"length")

countsDF = pairDF.groupBy("letter").agg(count("letter").alias("totalCount"),
sum("length").alias("totalLength"))

avgDF = countsDF.withColumn("ratio",
countsDF.totalLength/countsDF.totalCount).select("letter", "ratio").orderBy("letter")

avgDF.write.format("csv").save(outputFolder)
```

Practice (Spark SQL)

- ❖ Problem 3: Given a collection of documents, compute the average length of words starting with each letter.

```
fileDF = spark.read.text(inputFile)

fileDF.selectExpr("explode(split(value, ' ')) as word").createOrReplaceTempView("words")

spark.sql("select * from words where length(word)>=1 and substr(word, 0, 1)>='a' and substr(word, 0, 1)<='z' ").createOrReplaceTempView("filteredwords")

spark.sql("select substr(word, 0, 1) as letter, length(word) as length from filteredwords").createOrReplaceTempView("pair")

spark.sql("select letter, sum(length) as totalLength, count(*) as totalCount from pair group by letter").createOrReplaceTempView("count")

avgDF = spark.sql("select letter, totalLength/totalCount as ratio from count order by letter")

avgDF.write.format("csv").save(outputFolder)
```

Pass a Function to RDD Operations

- ❖ For a given text file, find the longest word from each line.

```
from pyspark import SparkContext, SparkConf
import sys

def findlongest(termList):
    maxTerm = ""
    for t in termList:
        if len(t) > len(maxTerm):
            maxTerm = t
    return maxTerm

class wordCount:
    def run(self, inputPath, outputPath):
        sc = SparkContext("local", "longest")
        fileRDD = sc.textFile(inputPath)
        wordsRDD = fileRDD.map(lambda line: line.lower().split())
        longestRDD = wordsRDD.map(lambda termList: findlongest(termList))
        longestRDD.saveAsTextFile(outputPath)
        sc.stop()

if __name__ == "__main__":
    if len(sys.argv) != 3:
        print("wrong inputs")
        sys.exit(-1)
    wordCount().run(sys.argv[1], sys.argv[2])
```

Define Your Own Function with UDF

- ❖ PySpark UDF (a.k.a User Defined Function) is the most useful feature of Spark SQL & DataFrame that is used to extend the PySpark build in capabilities.
- ❖ PySpark UDF's are similar to UDF on traditional databases. In PySpark, you create a function in a Python syntax and wrap it with PySpark SQL `udf()` or register it as `udf` and use it on DataFrame and SQL respectively.
- ❖ <https://sparkbyexamples.com/pyspark/pyspark-udf-user-defined-function>
- ❖ <https://spark.apache.org/docs/latest/api/python/reference/pyspark.sql/api/pyspark.sql.functions.udf.html#pyspark.sql.functions.udf>

Pass a Function to DataFrame Operations

- ❖ For a given text file, find the longest word from each line.

```
from pyspark.sql.session import SparkSession
from pyspark.sql.functions import *
import sys

def findlongest(termList):
    maxTerm = ""
    for t in termList:
        if len(t) > len(maxTerm):
            maxTerm = t
    return maxTerm

class WordCount:
    def run(self, inputPath, outputPath):
        spark = SparkSession.builder.master("local").appName("longest").getOrCreate()
        fileDF = spark.read.text(inputPath)
        wordsDF = fileDF.select(split(fileDF.value, ' ').alias('termList'))
        longestUDF = udf(lambda termList: findlongest(termList))
        resDF = wordsDF.withColumn('longest', longestUDF('termList')).select('longest')
        resDF.write.text(outputPath)
        spark.stop()

if __name__ == "__main__":
    if len(sys.argv) != 3:
        print("Wrong inputs")
        sys.exit(-1)
    wordCount().run(sys.argv[1], sys.argv[2])
```


References

- ❖ <http://spark.apache.org/docs/latest/index.html>
- ❖ Spark SQL guide: <http://spark.apache.org/docs/latest/sql-programming-guide.html>
- ❖ <https://spark.apache.org/docs/3.3.0/api/scala/org/apache/spark/index.html>
- ❖ https://spark.apache.org/docs/3.3.0/api/python/getting_started/index.html
- ❖ [Learning Spark](#). 2nd edition
- ❖ Spark SQL Functions: <https://sparkbyexamples.com/spark/spark-sql-functions/>

End of Chapter 5.2