



Graph Definitions

Graphs

6/188

Many applications require

- a collection of *items* (i.e. a set)
- *relationships*/connections between items

Examples:

- maps: items are cities, connections are roads
- web: items are pages, connections are hyperlinks

Collection types you're familiar with

- arrays and lists ... linear sequence of items

Graphs are more general ... allow arbitrary connections


... Graphs

7/188

A graph  $G = (V,E)$

- $V$  is a set of *vertices*
- $E$  is a set of *edges* (subset of  $V \times V$ )

Example:



... Graphs

8/188

A real example: Australian road distances

Distance	Adelaide	Brisbane	Canberra	Darwin	Melbourne	Perth	Sydney
Adelaide	-	2055	-	3051	732	2716	-
Brisbane	2055	-	-	3429	1671	-	982
Canberra	-	-	-	-	658	-	309
Darwin	3051	3429	-	-	-	4049	-


Melbourne	732	1671	658	-	-	-	873
Perth	2716	-	-	4049	-	-	-
Sydney	-	982	309	-	873	-	-

Notes: vertices are cities, edges are distance between cities, symmetric

... Graphs

9/188

Alternative representation of above:



... Graphs

10/188

Questions we might ask about a graph:

- is there a way to get from item A to item B?
- what is the best way to get from A to B?
- which items are connected?

Graph algorithms are generally more complex than linked list ones:

- no implicit order of items
- graphs may contain cycles
- concrete representation is less obvious
- algorithm complexity depends on connection complexity

Properties of Graphs

11/188

Terminology:  $|V|$  and  $|E|$  (cardinality) normally written just as  $V$  and  $E$ .

A graph with  $V$  vertices has at most  $V(V-1)/2$  edges.

The ratio  $E:V$  can vary considerably.

- if  $E$  is closer to  $V^2$ , the graph is *dense*
- if  $E$  is closer to  $V$ , the graph is *sparse*
  - Example: web pages and hyperlinks, intersections and roads on street map

Knowing whether a graph is sparse or dense is important

- may affect choice of data structures to represent graph

- may affect choice of algorithms to process graph

## Exercise #1: Number of Edges

12/188

The edges in a graph represent pairs of connected vertices. A graph with  $V$  has  $V^2$  such pairs.

Consider  $V = \{1,2,3,4,5\}$  with all possible pairs:

$$E = \{ (1,1), (1,2), (1,3), (1,4), (1,5), (2,1), (2,2), \dots, (4,5), (5,5) \}$$

Why do we say that the maximum #edges is  $V(V-1)/2$ ?

... because

- $(v,w)$  and  $(w,v)$  denote the same edge (in an undirected graph)
- we do not consider loops  $(v,v)$  (in undirected graphs)

## Graph Terminology

14/188

For an edge  $e$  that connects vertices  $v$  and  $w$

- $v$  and  $w$  are *adjacent* (neighbours)
- $e$  is *incident* on both  $v$  and  $w$

*Degree* of a vertex  $v$

- number of edges incident on  $v$

Synonyms:

- vertex = node, edge = arc = link (NB: some people use arc for *directed* edges)

## ... Graph Terminology

15/188

*Path*: a sequence of vertices where

- each vertex has an edge to its predecessor

*Simple path*: a path where


- all vertices and edges are different

*Cycle*: a path

- that is simple except last vertex = first vertex

*Length* of path or cycle:

- #edges



## ... Graph Terminology

16/188

*Connected graph*

- there is a *path* from each vertex to every other vertex
- if a graph is not connected, it has  $\geq 2$  *connected components*

*Complete graph  $K_V$*

- there is an *edge* from each vertex to every other vertex
- in a complete graph,  $E = V(V-1)/2$



## ... Graph Terminology

17/188

*Tree*: connected (sub)graph with no cycles

*Spanning tree*: tree containing all vertices

*n-Clique*: complete subgraph on  $n$  nodes

Consider the following single graph:



This graph has 26 vertices, 33 edges, and 4 connected components

Note: The entire graph has no spanning tree; what is shown in green is a spanning tree of the third connected component

## ... Graph Terminology

18/188

A *spanning tree* of connected graph  $G = (V,E)$

- is a subgraph of  $G$  containing all of  $V$
- and is a single tree (connected, no cycles)


A *spanning forest* of non-connected graph  $G = (V,E)$

- is a subgraph of  $G$  containing all of  $V$

- and is a set of trees (not connected, no cycles),
  - with one tree for each *connected component*

## Exercise #2: Graph Terminology

19/188



1. How many edges need to be removed to obtain a spanning tree?
2. How many different spanning trees?

1.  $\frac{5 \cdot 4}{2} - 2 = 8$  spanning trees (no spanning tree if we remove  $\{e1, e2\}$  or  $\{e3, e4\}$ )

## ... Graph Terminology

21/188


### Undirected graph

- $edge(u, v) = edge(v, u)$ , no self-loops (i.e. no  $edge(v, v)$ )

### Directed graph

- $edge(u, v) \neq edge(v, u)$ , can have self-loops (i.e.  $edge(v, v)$ )

Example:



### Weighted graph

- each edge has an associated value (weight)
- e.g. road map (weights on edges are distances between cities)

Other types of graphs ...


### Multi-graph

- allow multiple edges between two vertices
- e.g. function call graph ( $f()$  calls  $g()$  in several places)

## Graph Data Structures

## Graph Representations

Four representations of the *same* graph:



We will discuss three different graph data structures:

1. Array of edges
2. Adjacency matrix
3. Adjacency list

## Array-of-edges Representation

24/188

Edges are represented as an array of Edge values (= pairs of vertices)

- disadvantage: deleting edges is slightly complex
- undirected: order of vertices in an Edge  $(v, w)$  doesn't matter



For simplicity, we always assume vertices to be numbered  $0 \dots V-1$

## ... Array-of-edges Representation

25/188

Graph initialisation

```
newGraph(V):
|   Input  number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate enough memory for g.edges[]
|   return g
```

How much is enough? ... No more than  $V(V-1)/2$  ... Much less in practice (sparse graph)

## ... Array-of-edges Representation

26/188

Edge insertion

```
insertEdge(g, (v, w)):
```

```
Input graph g, edge (v,w) // assumption: (v,w) not in g

g.edges[g.nE]=(v,w)
g.nE=g.nE+1

Edge removal

removeEdge(g,(v,w)):
    Input graph g, edge (v,w) // assumption: (v,w) in g

    i=0
    while (v,w)≠g.edges[i] do
        i=i+1
    end while
    g.edges[i]=g.edges[g.nE-1] // replace (v,w) by last edge in array
    g.nE=g.nE-1
```

Cost Analysis

27/188

Storage cost:  $O(E)$

Cost of operations:

- initialisation:  $O(1)$
- insert edge:  $O(1)$  (assuming edge array has space)
- find/delete edge:  $O(E)$  (need to find edge in edge array)

If array is full on insert

- allocate space for a bigger array, copy edges across  $\Rightarrow O(E)$

If we maintain edges in order

- use binary search to insert/find edge  $\Rightarrow O(\log E)$   
(requires binary search *tree* of edges  $\rightarrow$  week 4)

Exercise #3: Array-of-edges Representation

28/188

Assuming an array-of-edges representation ...

Write an algorithm to output all edges of the graph


```
show(g):
    Input graph g

    for all i=0 to g.nE-1 do
        print g.edges[i]
    end for

Time complexity:  $O(E)$ 
```

## Adjacency Matrix Representation

Edges represented by a  $V \times V$  matrix



- easily implemented as 2-dimensional array
- can represent graphs, digraphs and weighted graphs
  - graphs: symmetric boolean matrix
  - digraphs: non-symmetric boolean matrix
  - weighted: non-symmetric matrix of weight values
- disadvantage: if few edges (sparse)  $\Rightarrow$  memory-inefficient

... Adjacency Matrix Representation

31/188

Graph initialisation

```
newGraph(V):
    Input number of nodes V
    Output new empty graph

    g.nV = V // #vertices (numbered 0..V-1)
    g.nE = 0 // #edges
    allocate memory for g.edges[][]
    for all i,j=0..V-1 do
        g.edges[i][j]=0 // false
    end for
    return g
```

... Adjacency Matrix Representation

32/188

Edge insertion

```
insertEdge(g,(v,w)):
    Input graph g, edge (v,w)

    if g.edges[v][w]=0 then // (v,w) not in graph
        g.edges[v][w]=1 // set to true
        g.edges[w][v]=1
        g.nE=g.nE+1
    end if
```

## Edge removal

```
removeEdge(g, (v,w)) :  
|   Input graph g, edge (v,w)  
|  
|   if g.edges[v][w]≠0 then // (v,w) in graph  
|       g.edges[v][w]=0      // set to false  
|       g.edges[w][v]=0  
|       g.nE=g.nE-1  
|   end if
```

## Exercise #4: Adjacency-matrix Representation

33/188

Assuming an adjacency matrix representation ...

Write an algorithm to output all edges of the graph (no duplicates!)

## ... Adjacency Matrix Representation

34/188

```
show(g):  
|   Input graph g  
|  
|   for all i=0 to g.nV-2 do  
|       for all j=i+1 to g.nV-1 do  
|           if g.edges[i][j] then  
|               print i-"-"  
|           end if  
|       end for  
|   end for
```

Time complexity:  $O(V^2)$

## Exercise #5:

35/188

Analyse storage cost and time complexity of adjacency matrix representation

Storage cost:  $O(V^2)$

If the graph is sparse, most storage is wasted.


Cost of operations:

- initialisation:  $O(V^2)$  (initialise  $V \times V$  matrix)
- insert edge:  $O(1)$  (set two cells in matrix)
- delete edge:  $O(1)$  (unset two cells in matrix)

## Adjacency List Representation


37/188

For each vertex, store linked list of adjacent vertices:



Undirected graph

A[0] = <1, 3>  
A[1] = <0, 3>  
A[2] = <3>  
A[3] = <0, 1, 2>



Directed graph

A[0] = <3>  
A[1] = <0, 3>  
A[2] = <>  
A[3] = <2>

- relatively easy to implement in languages like C
- can represent graphs and digraphs
- memory efficient if  $E:V$  relatively small

- disadvantage: one graph has many possible representations (unless lists are ordered by same criterion e.g. ascending)

## ... Adjacency List Representation

38/188

Graph initialisation

```
newGraph(V):  
|   Input number of nodes V  
|   Output new empty graph  
|  
|   g.nV = V // #vertices (numbered 0..V-1)  
|   g.nE = 0 // #edges  
|   allocate memory for g.edges[]  
|   for all i=0..V-1 do  
|       g.edges[i]=NULL // empty list  
|   end for  
|   return g
```

## ... Adjacency List Representation

39/188

Edge insertion:

```
insertEdge(g, (v,w)) :  
|   Input graph g, edge (v,w)  
|  
|   insertLL(g.edges[v],w)  
|   insertLL(g.edges[w],v)  
|   g.nE=g.nE+1
```

Edge removal:

```
removeEdge(g, (v,w)) :  
|   Input graph g, edge (v,w)
```

```
deleteLL(g.edges[v],w)
deleteLL(g.edges[w],v)
g.nE=g.nE-1
```

Exercise #6:

40/188

Analyse storage cost and time complexity of adjacency list representation

Storage cost:  $O(V+E)$  ( $V$  list pointers, total of  $2\cdot E$  list elements)

- the larger of  $V,E$  determines the complexity

Cost of operations:

- initialisation:  $O(V)$  (initialise  $V$  lists)
- insert edge:  $O(I)$  (insert one vertex into list)
  - if you don't check for duplicates
- find/delete edge:  $O(V)$  (need to find vertex in list)

Comparison of Graph Representations

42/188

	array of edges	adjacency matrix	adjacency list
space usage	$E$	$V^2$	$V+E$
initialise	$I$	$V^2$	$V$
insert edge	$I$	$I$	$I$
find/delete edge	$E$	$I$	$V$

Other operations:

	array of edges	adjacency matrix	adjacency list
disconnected(v)?	$E$	$V$	$I$
isPath(x,y)?	$E\cdot\log V$	$V^2$	$V+E$
copy graph	$E$	$V^2$	$V+E$
destroy graph	$I$	$V$	$V+E$

Graph Abstract Data Type

Graph ADT

44/188

Data:

- set of edges, set of vertices

Operations:

- building: create graph, add edge
- deleting: remove edge, drop whole graph
- scanning: check if graph contains a given edge

Things to note:

- set of vertices is fixed when graph initialised
- we treat vertices as `ints`, but could be arbitrary `Items` (e.g. individual profiles on a social network)

... Graph ADT

45/188

Graph ADT interface `Graph.h`

```
// graph representation is hidden
typedef struct GraphRep *Graph;

// vertices denoted by integers 0..N-1
typedef int Vertex;

// edges are pairs of vertices (end-points)
typedef struct Edge { Vertex v; Vertex w; } Edge;

// operations on graphs
Graph newGraph(int V); // new graph with V vertices
int numOfVertices(Graph); // get number of vertices in a graph
void insertEdge(Graph, Edge);
void removeEdge(Graph, Edge);
bool adjacent(Graph, Vertex, Vertex); /* is there an edge
between two vertices */


void showGraph(Graph); // print all edges in a graph
void freeGraph(Graph);
```

Graph ADT (Array of Edges)

46/188

Implementation of GraphRep (array-of-edges representation)

```
typedef struct GraphRep {
    Edge *edges; // array of edges
    int nV; // #vertices (numbered 0..nV-1)
    int nE; // #edges
    int n; // size of edge array
} GraphRep;
```




Graph ADT (Adjacency Matrix)

47/188

Implementation of GraphRep (adjacency-matrix representation)

```
typedef struct GraphRep {
    int **edges; // adjacency matrix
    int  nV;     // #vertices
    int  nE;     // #edges
} GraphRep;
```



... Graph ADT (Adjacency Matrix)

48/188

Implementation of graph initialisation (adjacency-matrix representation)

```
Graph newGraph(int V) {
    assert(V >= 0);
    int i;

    Graph g = malloc(sizeof(GraphRep));    assert(g != NULL);
    g->nV = V;  g->nE = 0;

    // allocate memory for each row
    g->edges = malloc(V * sizeof(int *));    assert(g->edges != NULL);
    // allocate memory for each column and initialise with 0
    for (i = 0; i < V; i++) {
        g->edges[i] = calloc(V, sizeof(int)); assert(g->edges[i] != NULL);
    }
    return g;
}
```

standard library function `calloc(size_t nelems, size_t nbytes)`

- allocates a memory block of size `nelems*nbytes`
- and sets all bytes in that block to `zero`

... Graph ADT (Adjacency Matrix)

49/188

Implementation of edge insertion/removal (adjacency-matrix representation)

```
// check if vertex is valid in a graph
bool validV(Graph g, Vertex v) {
    return (g != NULL && v >= 0 && v < g->nV);
}

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (!g->edges[e.v][e.w]) { // edge e not in graph
        g->edges[e.v][e.w] = 1;
        g->edges[e.w][e.v] = 1;
        g->nE++;
    }
}
```

```
    }
}

void removeEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (g->edges[e.v][e.w]) { // edge e in graph
        g->edges[e.v][e.w] = 0;
        g->edges[e.w][e.v] = 0;
        g->nE--;
    }
}
```

Exercise #7: Checking Neighbours

50/188

Assuming an adjacency-matrix representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```
bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g,x) && validV(g,y));

    return (g->edges[x][y] != 0);
}
```


Graph ADT (Adjacency List)

52/188

Implementation of GraphRep (adjacency-list representation)

```
typedef struct GraphRep {
    Node **edges; // array of lists
    int  nV;      // #vertices
    int  nE;      // #edges
} GraphRep;
```

```
typedef struct Node { // linked list node
    Vertex  v;
    struct Node *next;
} Node;
```



Graph Traversal

# Finding a Path

Questions on paths:

- is there a path between two given vertices (*src,dest*)?
- what is the sequence of vertices from *src* to *dest*?

Approach to solving problem:


- examine vertices adjacent to *src*
- if any of them is *dest*, then done
- otherwise try vertices two edges from *src*
- repeat looking further and further from *src*

Two strategies for graph traversal/search: *depth-first*, *breadth-first*

- DFS follows one path to completion before considering others
- BFS "fans-out" from the starting vertex ("spreading" subgraph)

## ... Finding a Path

Comparison of BFS/DFS search for checking if there is a path from *a* to *h* ...



Both approaches ignore some edges by remembering previously visited vertices.

# Depth-first Search

Depth-first search can be described recursively as

**depthFirst(G,v) :**

1. mark *v* as visited
2. for each (*v,w*) $\in$ *edges*(*G*) do  
    if *w* has not been visited then  
        **depthFirst(w)**

The recursion induces *backtracking*

## ... Depth-first Search

Recursive DFS path checking

**hasPath(G,src,dest) :**

54/188


**Input** graph *G*, vertices *src,dest*  
**Output** true if there is a path from *src* to *dest* in *G*,  
false otherwise

mark all vertices in *G* as unvisited  
**return** **dfsPathCheck(G,src,dest)**

```
dfsPathCheck(G,v,dest):  
    mark v as visited  
    if v=dest then           // found dest  
        return true  
    else  
        for all (v,w) $\in$ edges(G) do  
            if w has not been visited then  
                if dfsPathCheck(G,w,dest) then  
                    return true // found path via w to dest  
                end if  
            end if  
        end for  
    end if  
    return false             // no path from v to dest
```

## Exercise #8: Depth-first Traversal (i)

Trace the execution of **dfsPathCheck(G,0,5)** on:



Consider neighbours in ascending order

Answer:

0 - 1 - 2 - 3 - 4 - 5

## ... Depth-first Search

Cost analysis:

- all vertices marked as unvisited, each vertex visited at most once  $\Rightarrow$  cost =  $O(V)$
- visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(E)$ 
  - assuming an adjacency list representation

*Time complexity of DFS:*  $O(V+E)$  (adjacency list representation)

- the *larger* of  $V,E$  determines the complexity

57/188

58/188

60/188



For *dense graphs* ...  $E \cong V^2 \Rightarrow O(V+E) = O(V^2)$   
For *sparse graphs* ...  $E \cong V \Rightarrow O(V+E) = O(V)$

... Depth-first Search

61/188

Note how different graph data structures affect cost:

- array-of-edges representation
  - visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(V \cdot E)$
  - cost of DFS:  $O(V \cdot E)$
- adjacency-matrix representation
  - visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(V^2)$
  - cost of DFS:  $O(V^2)$

... Depth-first Search

62/188

Knowing whether a path exists can be useful

Knowing what the path is even more useful  
 $\Rightarrow$  record the previously visited node as we search through the graph (so that we can then trace path through graph)

Make use of global variable:

- `visited[ ]` ... array to store previously visited node, for each node being visited

... Depth-first Search

63/188

`visited[ ]` // store previously visited node, for each vertex  $0..nV-1$

```
findPath(G,src,dest):
  Input graph G, vertices src,dest

  for all vertices v∈G do
    visited[v]=-1
  end for
  visited[src]=src // starting node of the path
  if dfsPathCheck(G,src,dest) then // show path in dest..src order
    v=dest
    while v≠src do
      print v "-"
      v=visited[v]
    end while
    print src
  end if


dfsPathCheck(G,v,dest):
  if v=dest then // found edge from v to dest
    return true
  else
    for all (v,w)∈Edges(G) do
      if visited[w]=-1 then
        visited[w]=v
        if dfsPathCheck(G,w,dest) then
```

```
      return true // found path via w to dest
    end if
  end if
end for
end if
return false // no path from v to dest
```

Exercise #9: Depth-first Traversal (ii)

64/188

Show the DFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	3	5	3	1	5	4	7	8
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-1-0

... Depth-first Search

66/188

DFS can also be described non-recursively (via a *stack*):

```
hasPath(G,src,dest):
  Input graph G, vertices src,dest
  Output true if there is a path from src to dest in G,
         false otherwise

  mark all vertices in G as unvisited
  push src onto new stack s
  found=false
  while not found and s is not empty do
    pop v from s
    mark v as visited
    if v=dest then
      found=true
    else
      for each (v,w)∈Edges(G) such that w has not been visited
        push w onto s
      end for
    end if
  end while
  return found
```


Uses standard stack operations (push, pop, check if empty)

Time complexity is the same:  $O(V+E)$

Exercise #10: Depth-first Traversal (iii)

67/188

Show how the stack evolves when executing findPathDFS(*g*,0,5) on:



Push neighbours in *descending* order ... so they get popped in ascending order

			4	5		
		3	5	5	5	
1	2	4	4	4	4	
4	4	4	4	4	4	
(empty)	→ 0	→ 5	→ 5	→ 5	→ 5	→ 5

Breadth-first Search

69/188

Basic approach to breadth-first search (BFS):

- visit and mark current vertex
- visit all neighbours of current vertex
- then consider neighbours of neighbours

Notes:

- tricky to describe recursively
- a minor variation on non-recursive DFS search works  
⇒ switch the *stack* for a *queue*

... Breadth-first Search

70/188

BFS algorithm (records visiting order, marks vertices as visited when put *on* queue):

visited[] // array of visiting orders, indexed by vertex 0..*nV*-1

```
findPathBFS(G,src,dest):
|   Input  graph G, vertices src,dest
|
|   for all vertices v∈G do
```


```
    visited[v]=-1
end for
enqueue src into new queue q
visited[src]=src
found=false
while not found and q is not empty do
|   dequeue v from q
|   if v=dest then
|       found=true
|   else
|       for each (v,w)∈edges(G) such that visited[w]=-1 do
|           enqueue w into q
|           visited[w]=v
|       end for
|   end if
end while
if found then
    display path in dest..src order
end if
```

Uses standard queue operations (enqueue, dequeue, check if empty)

Exercise #11: Breadth-first Traversal

71/188

Show the state of the queue and the BFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



visited	?	?	?	?	?	?	?	?	?
	[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

Consider neighbours in ascending order

Queue (front to the left):

(empty)	→ 0	→ 1 2	→ 2	→ 5	→ 3 4 6	→ 4 6 7	→ 6 7	→ 7
		5	5	3	7	8	8	8

0	0	0	2	5	0	5	5	3	-1
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-0

... Breadth-first Search

73/188

*Time complexity of BFS:  $O(V+E)$*  (adjacency list representation, same as DFS)

BFS finds a "shortest" path

- based on minimum # edges between *src* and *dest*.
- stops with first-found path, if there are multiple ones

In many applications, edges are weighted and we want path

- based on minimum sum-of-weights along path *src* .. *dest*


We discuss weighted/directed graphs later.

## Other DFS Examples

74/188

Other problems to solve via DFS graph search

- checking for the existence of a cycle
- determining which connected component each vertex is in



## Exercise #12: Buggy Cycle Check

75/188

A graph has a *cycle* if

- it has a path of length  $> 1$
- with start vertex *src* = end vertex *dest*
- and without using any edge more than once

We are not required to give the path, just indicate its presence.

The following DFS cycle check has two bugs. Find them.

```
hasCycle(G):
|   Input  graph G
|   Output true if G has a cycle, false otherwise
|
|   choose any vertex  $v \in G$ 
|   return dfsCycleCheck(G,v)
|
dfsCycleCheck(G,v):
|   mark v as visited
|   for each  $(v,w) \in \text{edges}(G)$  do
|   |   if w has been visited then    // found cycle
|   |   |   return true
|   |   else if dfsCycleCheck(G,w) then
|   |   |   return true
|   |   end if
|   end for
|   return false                    // no cycle at v
```

1. Only one connected component is checked.
2. The loop

**for each**  $(v,w) \in \text{edges}(G)$  **do**

should exclude the neighbour of *v* from which you just came, so as to prevent a single edge *w-v* from being classified as a cycle.

## Computing Connected Components

77/188

Problems:

- how many connected subgraphs are there?
- are two vertices in the same connected subgraph?

Both of the above can be solved if we can

- build an array, one element for each vertex *V*
- indicating which connected component *V* is in
- componentOf[ ] ... array [0..nV-1] of component IDs

## ... Computing Connected Components

78/188

Algorithm to assign vertices to connected components:

```
components(G):
|   Input  graph G
|
|   for all vertices  $v \in G$  do
|   |   componentOf[v] = -1
|   end for
|   compID = 0
|   for all vertices  $v \in G$  do
|   |   if componentOf[v] = -1 then
|   |   |   dfsComponents(G,v,compID)
|   |   |   compID = compID + 1
|   |   end if
|   end for
|
dfsComponents(G,v,id):
|   componentOf[v] = id
|   for all vertices w adjacent to v do
|   |   if componentOf[w] = -1 then
|   |   |   dfsComponents(G,w,id)
|   |   end if
|   end for
```


## Exercise #13: Connected components

79/188

Trace the execution of the algorithm

1. on the graph shown below

2. on the same graph but with the dotted edges added



Consider neighbours in ascending order

1.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	-1	0	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
0	0	0	1	-1	-1	-1	-1
...							
0	0	0	1	1	2	2	2

2.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	0	-1	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
...							
0	0	0	0	0	1	1	1

## Hamiltonian and Euler Paths

### Hamiltonian Path and Circuit

Hamiltonian path problem:

- find a path connecting two vertices  $v,w$  in graph  $G$
- such that the path includes each *vertex* exactly once

If  $v = w$ , then we have a *Hamiltonian circuit*

Simple to state, but difficult to solve (*NP*-complete)


Many real-world applications require you to visit all vertices of a graph:

- Travelling salesman
- Bus routes
- ...

Problem named after Irish mathematician, physicist and astronomer Sir William Rowan Hamilton (1805 — 1865)

### ... Hamiltonian Path and Circuit

Graph and two possible Hamiltonian paths:



### ... Hamiltonian Path and Circuit

Approach:

- generate all possible simple paths (using e.g. DFS)
- keep a counter of vertices visited in current path
- stop when find a path containing  $V$  vertices

Can be expressed via a recursive DFS algorithm

- similar to simple path finding approach, except
  - keeps track of path length; succeeds if  $\text{length} = v-1$  (length =  $v$  for circuit)
  - resets "visited" marker after unsuccessful path

### ... Hamiltonian Path and Circuit

Algorithm for finding Hamiltonian path:

visited[] // array [0.. $nV-1$ ] to keep track of visited vertices

```
hasHamiltonianPath(G,src,dest):
  for all vertices v in G do
    visited[v]=false
  end for
  return hamiltonR(G,src,dest,#vertices(G)-1)

hamiltonR(G,v,dest,d):
  Input G      graph
         v      current vertex considered
         dest   destination vertex
         d      distance "remaining" until path found

  if v=dest then
    if d=0 then return true else return false
  else
    mark v as visited
    for each neighbour w of v in G do
      if w has not been visited then
        if hamiltonR(G,w,dest,d-1) then
```

```


    return true
  end if
end if
end for
end if
mark v as unvisited      // reset visited mark
return false

```

### Exercise #14: Hamiltonian Path

86/188

Trace the execution of the algorithm when searching for a Hamiltonian path from 1 to 6:



Consider neighbours in ascending order

1-0-2-3-4-5-6	d≠0
1-0-2-3-4-5-7-8-9	no unvisited neighbour
1-0-2-3-4-5-7-9-8	no unvisited neighbour
1-0-2-3-4-7-5-6	d≠0
1-0-2-3-4-7-8-9	no unvisited neighbour
1-0-2-3-4-7-9-8	no unvisited neighbour
1-0-2-3-4-8-7-5-6	d≠0
1-0-2-3-4-8-7-9	no unvisited neighbour
1-0-2-3-4-8-9-7-5-6	✓


Repeat on your own with `src=0` and `dest=6`

### ... Hamiltonian Path and Circuit

88/188

Analysis: worst case requires  $(V-1)!$  paths to be examined

Consider a graph with isolated vertex and the rest fully-connected



Checking `hasHamiltonianPath(g,x,0)` for any  $x$

- requires us to consider every possible path
- e.g 1-2-3-4, 1-2-4-3, 1-3-2-4, 1-3-4-2, 1-4-2-3, ...
- starting from any  $x$ , there are  $3!$  paths  $\Rightarrow 4!$  total paths
- there is no path of length 5 in these  $(V-1)!$  possibilities

There is no known polynomial algorithm for this task ( $NP$ -complete)

Note, however, that the above case could be solved in constant time if we had a fast check for 0 and  $x$  being in the same connected component


## Euler Path and Circuit

89/188


*Euler path* problem:

- find a path connecting two vertices  $v,w$  in graph  $G$
- such that the path includes each *edge* exactly once  
(note: the path does not have to be simple  $\Rightarrow$  can visit vertices more than once)

If  $v = w$ , then we have an *Euler circuit*



Euler Path: 4-3-1-5-2-1



Euler Circuit: 1-2-5-4-3-1

Many real-world applications require you to visit all edges of a graph:

- Postman
- Garbage pickup
- ...

Problem named after Swiss mathematician, physicist, astronomer, logician and engineer Leonhard Euler (1707 — 1783)

### ... Euler Path and Circuit

90/188

One possible "brute-force" approach:

- check for each path if it's an Euler path
- would result in factorial time performance

Can develop a better algorithm by exploiting:


*Theorem.* A graph has an Euler circuit if and only if it is connected and all vertices have even degree


*Theorem.* A graph has a non-circuitous Euler path if and only if it is connected and exactly two vertices have odd degree

### Exercise #15: Euler Paths and Circuits

91/188

Which of these two graphs have an Euler path? an Euler circuit?





No Euler circuit

Only the second graph has an Euler path, e.g. 2-0-1-3-2-4-5

### ... Euler Path and Circuit

93/188

Assume the existence of  $\text{degree}(g, v)$  (degree of a vertex, cf. problem set 1 exercise 2 this week)

Algorithm to check whether a graph has an Euler path:

```
hasEulerPath(G, src, dest):
  Input  graph G, vertices src, dest
  Output true if G has Euler path from src to dest
         false otherwise

  if src ≠ dest then           // non-circuitous path
    if degree(G, src) or degree(G, dest) is even then
      return false
    end if
  else if degree(G, src) is odd then // circuit
    return false
  end if
  for all vertices v ∈ G do
    if v ≠ src and v ≠ dest and degree(G, v) is odd then
      return false
    end if
  end for
  return true
```

### ... Euler Path and Circuit

94/188

Analysis of hasEulerPath algorithm:

- assume that connectivity is already checked
- assume that degree is available via  $O(1)$  lookup
- single loop over all vertices  $\Rightarrow O(V)$

If degree requires iteration over vertices

- cost to compute degree of a single vertex is  $O(V)$
- overall cost is  $O(V^2)$

$\Rightarrow$  problem tractable, even for large graphs (unlike Hamiltonian path problem)

For the keen: Linear-time (in the number of edges,  $E$ ) algorithm to compute an Euler path described in [Sedgewick] Ch.17.7

## Directed Graphs

### Directed Graphs (Digraphs)


102/188

In our previous discussion of graphs:

- an edge indicates a relationship between two vertices
- an edge indicates nothing more than a relationship

In many real-world applications of graphs:

- edges are directional ( $v \rightarrow w \neq w \rightarrow v$ )




- edges have a *weight* (cost to go from  $v \rightarrow w$ )

### ... Directed Graphs (Digraphs)

103/188

Example digraph and adjacency matrix representation:



Unidirectional  $\Rightarrow$  symmetric matrix

Directional  $\Rightarrow$  non-symmetric matrix

Maximum #edges in a digraph with  $V$  vertices:  $V^2$

### ... Directed Graphs (Digraphs)

104/188

Terminology for digraphs ...

*Directed path*: sequence of  $n \geq 2$  vertices  $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_n$

- where  $(v_i, v_{i+1}) \in \text{edges}(G)$  for all  $v_i, v_{i+1}$  in sequence
- if  $v_1 = v_n$ , we have a *directed cycle*

*Reachability*:  $w$  is reachable from  $v$  if  $\exists$  directed path  $v, \dots, w$

## Digraph Applications

105/188

Potential application areas:

Domain	Vertex	Edge
Web	web page	hyperlink
scheduling	task	precedence
chess	board position	legal move
science	journal article	citation
dynamic data	malloc'd object	pointer
programs	function	function call
make	file	dependency

## ... Digraph Applications

106/188

Problems to solve on digraphs:

- is there a directed path from  $s$  to  $t$ ? (transitive closure)
- what is the shortest path from  $s$  to  $t$ ? (shortest path)
- are all vertices mutually reachable? (strong connectivity)
- how to organise a set of tasks? (topological sort)
- which web pages are "important"? (PageRank)
- how to build a web crawler? (graph traversal)


## Digraph Representation

107/188

Similar set of choices as for undirectional graphs:

- array of edges (directed)
- vertex-indexed adjacency matrix (non-symmetric)
- vertex-indexed adjacency lists

$V$  vertices identified by  $0 \dots V-1$



## Reachability

### Transitive Closure

109/188

Given a digraph  $G$  it is potentially useful to know

- is vertex  $t$  reachable from vertex  $s$ ?

Example applications:

- can I complete a schedule from the current state?
- is a malloc'd object being referenced by any pointer?

How to compute transitive closure?

### ... Transitive Closure

110/188

One possibility:

- implement it via `hasPath( $G, s, t$ )` (itself implemented by DFS or BFS algorithm)
- feasible if `reachable( $G, s, t$ )` is infrequent operation

What if we have an algorithm that frequently needs to check reachability?

Would be very convenient/efficient to have:


```
reachable( $G, s, t$ ):
|   return  $G.tc[s][t]$     // transitive closure matrix
```

Of course, if  $V$  is very large, then this is not feasible.


### Exercise #16: Transitive Closure Matrix

111/188

Which reachable  $s \dots t$  exist in the following graph?



Transitive closure of example graph:



### ... Transitive Closure

113/188

**Goal:** produce a matrix of reachability values

- if  $tc[s][t]$  is 1, then  $t$  is reachable from  $s$
- if  $tc[s][t]$  is 0, then  $t$  is not reachable from  $s$

So, how to create this matrix?

Observation:

$\forall i,s,t \in \text{vertices}(G):$   
 $(s,i) \in \text{edges}(G) \text{ and } (i,t) \in \text{edges}(G) \Rightarrow tc[s][t] = 1$

$tc[s][t]=1$  if there is a path from  $s$  to  $t$  via some  $i \quad (s \rightarrow i \rightarrow t)$

... Transitive Closure

114/188

If we implement the above as:

```
make tc[][] a copy of edges[][]
for all i∈vertices(G) do
  for all s∈vertices(G) do
    for all t∈vertices(G) do
      if tc[s][i]=1 and tc[i][t]=1 then
        tc[s][t]=1
      end if
    end for
  end for
end for
```

then we get an algorithm to convert edges into a tc

This is known as *Warshall's algorithm*

... Transitive Closure

115/188

How it works ...

After iteration 1,  $tc[s][t]$  is 1 if

- either  $s \rightarrow t$  exists or  $s \rightarrow 0 \rightarrow t$  exists

After iteration 2,  $tc[s][t]$  is 1 if any of the following exist

- $s \rightarrow t$  or  $s \rightarrow 0 \rightarrow t$  or  $s \rightarrow 1 \rightarrow t$  or  $s \rightarrow 0 \rightarrow 1 \rightarrow t$  or  $s \rightarrow 1 \rightarrow 0 \rightarrow t$


Etc. ... so after the  $V^{th}$  iteration,  $tc[s][t]$  is 1 if

- there is any directed path in the graph from  $s$  to  $t$

Exercise #17: Transitive Closure

116/188

Trace Warshall's algorithm on the following graph:



1<sup>st</sup> iteration i=0:

tc	[0]	[1]	[2]	[3]
----	-----	-----	-----	-----

[0]	0	1	1	1
[1]	1	1	1	1
[2]	0	1	0	0
[3]	0	0	0	0

2<sup>nd</sup> iteration i=1:

tc	[0]	[1]	[2]	[3]
[0]	1	1	1	1
[1]	1	1	1	1
[2]	1	1	1	1
[3]	0	0	0	0

3<sup>rd</sup> iteration i=2: unchanged

4<sup>th</sup> iteration i=3: unchanged

... Transitive Closure

118/188

Cost analysis:

- storage: additional  $V^2$  items (each item may be 1 bit)
- computation of transitive closure:  $O(V^3)$
- computation of `reachable()`:  $O(I)$  after having generated `tc[][]`

Amortisation: would need many calls to `reachable()` to justify other costs

Alternative: use DFS in each call to `reachable()`

Cost analysis:

- storage: cost of stack and set ("visited") during `reachable()`
- computation of `reachable()`: cost of DFS =  $O(V^2)$  (for adjacency matrix)

Digraph Traversal

119/188

Same algorithms as for undirected graphs:

depthFirst(v):

1. mark  $v$  as visited
2. for each  $(v,w) \in \text{edges}(G)$  do  
if  $w$  has not been visited then  
**depthFirst(w)**

breadth-first(v):

1. enqueue  $v$
2. while queue not empty do  
dequeue  $v$   
if  $v$  not already visited then



mark  $v$  as visited  
enqueue each vertex  $w$  adjacent to  $v$

## Example: Web Crawling

120/188

**Goal:** visit every page on the web

**Solution:** breadth-first search with "implicit" graph

```
webCrawl(startingURL):
    mark startingURL as alreadySeen
    enqueue(Q, startingURL)
    while Q is not empty do
        nextPage=dequeue(Q)
        visit nextPage
        for each hyperLink on nextPage do
            if hyperLink not alreadySeen then
                mark hyperLink as alreadySeen
                enqueue(Q, hyperLink)
            end if
        end for
    end while
```

visit scans page and collects e.g. keywords and links

## Weighted Graphs

## Weighted Graphs

122/188

Graphs so far have considered

- edge = an association between two vertices/nodes
- may be a precedence in the association (directed)

Some applications require us to consider


- a *cost* or *weight* of an association
- modelled by assigning values to edges (e.g. positive reals)

Weights can be used in both directed and undirected graphs.

### ... Weighted Graphs

123/188

Example: major airline flight routes in Australia



Representation: edge = direct flight; weight = approx flying time (hours)

### ... Weighted Graphs

124/188

Weights lead to minimisation-type questions, e.g.

1. Cheapest way to connect all vertices?

- a.k.a. *minimum spanning tree* problem
- assumes: edges are weighted and undirected

2. Cheapest way to get from  $A$  to  $B$ ?

- a.k.a. *shortest path* problem
- assumes: edge weights positive, directed or undirected

### Exercise #18: Implementing a Route Finder

125/188

If we represent a street map as a graph

- what are the vertices?
- what are the edges?
- are edges directional?
- what are the weights?
- are the weights fixed?

## Weighted Graph Representation

126/188

Weights can easily be added to:


- adjacency matrix representation (0/1  $\rightarrow$  int or float)
- adjacency lists representation (add int/float to list node)

Both representations work whether edges are directed or not.

### ... Weighted Graph Representation

127/188

Adjacency matrix representation with weights:




Note: need distinguished value to indicate "no edge".

## ... Weighted Graph Representation

128/188

Adjacency lists representation with weights:



Note: if undirected, each edge appears twice with same weight

## ... Weighted Graph Representation

129/188

Sample adjacency matrix implementation in C requires minimal changes to previous Graph ADT:

### WGraph.h

```
// edges are pairs of vertices (end-points) plus positive weight
typedef struct Edge {
    Vertex v;
    Vertex w;
    int weight;
} Edge;

// returns weight, or 0 if vertices not adjacent
int adjacent(Graph, Vertex, Vertex);
```

## ... Weighted Graph Representation

130/188

### WGraph.c

```
typedef struct GraphRep {
    int **edges; // adjacency matrix storing positive weights
                // 0 if nodes not adjacent
    int nV;      // #vertices
    int nE;      // #edges
} GraphRep;

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));
    if (g->edges[e.v][e.w] == 0) { // edge e not in graph
        g->edges[e.v][e.w] = e.weight;
        g->nE++;
    }
```

```
}
}


int adjacent(Graph g, Vertex v, Vertex w) {
    assert(g != NULL && validV(g,v) && validV(g,w));
    return g->edges[v][w];
}
```

## Minimum Spanning Trees

### Exercise #19: Minimising Wires in Circuits

132/188

Electronic circuit designs often need to make the pins of several components electrically equivalent by wiring them together.



To interconnect a set of  $n$  pins we can use an arrangement of  $n-1$  wires each connecting two pins.

What kind of algorithm would ...

- help us find the arrangement with the least amount of wire?

## Minimum Spanning Trees

133/188

Reminder: *Spanning tree ST* of graph  $G=(V,E)$

- *spanning* = all vertices, *tree* = no cycles
  - *ST* is a subgraph of  $G$  ( $G'=(V,E')$  where  $E' \subseteq E$ )
  - *ST* is *connected* and *acyclic*

*Minimum spanning tree MST* of graph  $G$

- *MST* is a spanning tree of  $G$
- sum of edge weights is no larger than any other ST

Applications: Computer networks, Electrical grids, Transportation networks ...


**Problem:** how to (efficiently) find MST for graph  $G$ ?

NB: MST may not be unique (e.g. all edges have same weight  $\Rightarrow$  every ST is MST)


... Minimum Spanning Trees

134/188

Example:



An MST ...



... Minimum Spanning Trees

135/188

Brute force solution:

```
findMST(G):
  Input graph G
  Output a minimum spanning tree of G

  bestCost=∞
  for all spanning trees t of G do
    if cost(t)<bestCost then
      bestTree=t
      bestCost=cost(t)
    end if
  end for
  return bestTree
```

Example of *generate-and-test* algorithm.

Not useful because #spanning trees is potentially large (e.g.  $n^{n-2}$  for a complete graph with  $n$  vertices)

... Minimum Spanning Trees

136/188

Simplifying assumption:

- edges in  $G$  are not directed (MST for digraphs is harder)

Kruskal's Algorithm

137/188

One approach to computing MST for graph  $G$  with  $V$  nodes:

1. start with empty MST
2. consider edges in increasing weight order
  - add edge if it does not form a cycle in MST
3. repeat until  $V-1$  edges are added


Critical operations:

- iterating over edges in weight order
- checking for cycles in a graph

... Kruskal's Algorithm

138/188


Execution trace of Kruskal's algorithm:




Exercise #20: Kruskal's Algorithm

139/188


Show how Kruskal's algorithm produces an MST on:




After 3<sup>rd</sup> iteration:




After 6<sup>th</sup> iteration:



After 7<sup>th</sup> iteration:



After 8<sup>th</sup> iteration ( $V-1=8$  edges added):



## ... Kruskal's Algorithm

141/188

Pseudocode:

```
KruskalMST(G):
  Input  graph G with n nodes
  Output a minimum spanning tree of G

  MST=empty graph
  sort edges(G) by weight
  for each e in sortedEdgeList do
    MST = MST U {e}
    if MST has a cycle then
      MST = MST \ {e}
    end if
    if MST has n-1 edges then
      return MST
    end if
  end for
```

## ... Kruskal's Algorithm

142/188

Time complexity analysis ...

- sorting edge list is  $O(E \cdot \log E)$
- min  $V$ , max  $E$  iterations over sorted edges
- on each iteration ...
  - getting next lowest cost edge is  $O(1)$
  - checking whether adding it forms a cycle: cost = ??
    - use DFS ... too expensive?
    - could use *Union-Find data structure* (see Sedgewick Ch.1) to maintain sets of connected components
      - ⇒ loop is  $O(E \cdot \log V)$
- overall complexity  $O(E \cdot \log E) = O(E \cdot \log V)$

## Exercise #21: Kruskal's Algorithm

143/188

Why is  $O(E \cdot \log E) = O(E \cdot \log V)$  in this case?

1. at most  $E = V^2$  edges  $\Rightarrow \log E = 2 \cdot \log V = O(\log V)$
2. if  $V > E+1 \Rightarrow$  can ignore all unconnected vertices

## Prim's Algorithm

145/188

Another approach to computing MST for graph  $G=(V,E)$ :

1. start from any vertex  $v$  and empty MST
2. choose edge not already in MST to add to MST
  - must be incident on a vertex  $s$  already connected to  $v$  in MST
  - must be incident on a vertex  $t$  not already connected to  $v$  in MST
  - must have minimal weight of all such edges
3. repeat until MST covers all vertices


Critical operations:

- checking for vertex being connected in a graph
- finding min weight edge in a set of edges

## ... Prim's Algorithm

146/188


Execution trace of Prim's algorithm (starting at  $s=0$ ):



## Exercise #22: Prim's Algorithm


147/188

Show how Prim's algorithm produces an MST on:




Start from vertex 0


After 1<sup>st</sup> iteration:




After 2<sup>nd</sup> iteration:




After 3<sup>rd</sup> iteration:



After 4<sup>th</sup> iteration:



After 8<sup>th</sup> iteration (all vertices covered):



## ... Prim's Algorithm

149/188

Pseudocode:

```
PrimMST(G):
  Input graph G with n>0 nodes
  Output a minimum spanning tree of G

  MST=empty graph
  usedV={0}
  unusedE=edges(g)
  while |usedV|<n do
    find e=(s,t,w) ∈ unusedE such that {
      s ∈ usedV, t ∉ usedV and w is min weight of all such edges
    }
    MST = MST ∪ {e}
    usedV = usedV ∪ {t}
    unusedE = unusedE \ {e}
  end while
  return MST
```

Critical operation: finding best edge

## ... Prim's Algorithm

150/188

Rough time complexity analysis ...

- $V$  iterations of outer loop
- find min edge with set of edges is  $O(E) \Rightarrow O(V \cdot E)$  overall

Using a *priority queue* ...

- $\Rightarrow O(E \cdot \log V)$  overall

## Sidetrack: Priority Queues

151/188

Some applications of queues require

- items processed in order of "priority"
- rather than in order of entry (FIFO — first in, first out)

*Priority Queues (PQueues)* provide this via:

- **join**: insert item into PQueue with an associated priority (replacing enqueue)
- **leave**: remove item with highest priority (replacing dequeue)

Time complexity for naive implementation of a PQueue containing  $N$  items ...

- $O(1)$  for join     $O(N)$  for leave

Most efficient implementation ("heap") ...

- $O(\log N)$  for join, leave ... more on this in week 4 (binary search trees)

## Other MST Algorithms

152/188

**Boruvka's algorithm** ... complexity  $O(E \cdot \log V)$

- the oldest MST algorithm
- start with  $V$  separate components
- join components using min cost links
- continue until only a single component

**Karger, Klein, and Tarjan** ... complexity  $O(E)$

- based on Boruvka, but non-deterministic
- randomly selects subset of edges to consider
- for the keen, here's [the paper](#) describing the algorithm

## Shortest Path

## Shortest Path

154/188

*Path* = sequence of edges in graph  $G$      $p = (v_0, v_1), (v_1, v_2), \dots, (v_{m-1}, v_m)$

$cost(path)$  = sum of edge weights along path

*Shortest path* between vertices  $s$  and  $t$

- a simple path  $p(s,t)$  where  $s = first(p)$ ,  $t = last(p)$
- no other simple path  $q(s,t)$  has  $cost(q) < cost(p)$

Assumptions: weighted digraph, no negative weights.

Finding shortest path between two given nodes known as *source-target* SP problem

Variations: *single-source* SP, *all-pairs* SP

Applications: navigation, routing in data networks, ...

Single-source Shortest Path (SSSP)


155/188

Given: weighted digraph  $G$ , source vertex  $s$

Result: shortest paths from  $s$  to all other vertices

- `dist[]`  $V$ -indexed array of cost of shortest path from  $s$
- `pred[]`  $V$ -indexed array of predecessor in shortest path from  $s$

Example:




Edge Relaxation

156/188

Assume: `dist[]` and `pred[]` as above (but containing data for shortest paths *discovered so far*)

`dist[v]` is length of shortest known path from  $s$  to  $v$   
`dist[w]` is length of shortest known path from  $s$  to  $w$

*Relaxation* updates data for  $w$  if we find a shorter path from  $s$  to  $w$ :



Relaxation along edge  $e=(v,w,weight)$ :

- if `dist[v]+weight < dist[w]` then  
update `dist[w]:=dist[v]+weight` and `pred[w]:=v`

Dijkstra's Algorithm

157/188

One approach to solving single-source shortest path problem ...

Data:  $G, s, dist[], pred[]$  and

- `vSet`: set of vertices whose shortest path from  $s$  is unknown

Algorithm:

`dist[]` // array of cost of shortest path from  $s$   
`pred[]` // array of predecessor in shortest path from  $s$


```
dijkstraSSSP(G,source):
    Input graph G, source node

    initialise dist[] to all ∞, except dist[source]=0
    initialise pred[] to all -1
    vSet=all vertices of G
    while vSet≠∅ do
        find s∈vSet with minimum dist[s]
        for each (s,t,w)∈edges(G) do
            relax along (s,t,w)
        end for
        vSet=vSet\{s}
    end while
```

Exercise #23: Dijkstra's Algorithm

158/188

Show how Dijkstra's algorithm runs on (source node = 0):



	[0]	[1]	[2]	[3]	[4]	[5]
dist	0	∞	∞	∞	∞	∞
pred	-	-	-	-	-	-

	[0]	[1]	[2]	[3]	[4]	[5]
dist	0	∞	∞	∞	∞	∞
pred	-	-	-	-	-	-

dist	0	14	9	7	∞	∞
pred	-	0	0	0	-	-

dist	0	14	9	7	∞	22
pred	-	0	0	0	-	3

dist	0	13	9	7	∞	12

pred	-	2	0	0	-	2
------	---	---	---	---	---	---

dist	0	13	9	7	20	12
pred	-	2	0	0	5	2

dist	0	13	9	7	18	12
pred	-	2	0	0	1	2

### ... Dijkstra's Algorithm

160/188

Why Dijkstra's algorithm is correct:

*Hypothesis.*

- (a) For visited  $s \dots dist[s]$  is shortest distance from source
- (b) For unvisited  $t \dots dist[t]$  is shortest distance from source via visited nodes

*Proof.*

Base case: no visited nodes,  $dist[source]=0$ ,  $dist[s]=\infty$  for all other nodes

Induction step:

1. If  $s$  is unvisited node with minimum  $dist[s]$ , then  $dist[s]$  is shortest distance from source to  $s$ :
  - if  $\exists$  shorter path via only visited nodes, then  $dist[s]$  would have been updated when processing the predecessor of  $s$  on this path
  - if  $\exists$  shorter path via an unvisited node  $u$ , then  $dist[u]<dist[s]$ , which is impossible if  $s$  has min distance of all unvisited nodes
2. This implies that (a) holds for  $s$  after processing  $s$
3. (b) still holds for all unvisited nodes  $t$  after processing  $s$ :
  - if  $\exists$  shorter path via  $s$  we would have just updated  $dist[t]$
  - if  $\exists$  shorter path without  $s$  we would have found it previously

### ... Dijkstra's Algorithm

161/188

Time complexity analysis ...

Each edge needs to be considered once  $\Rightarrow O(E)$ .

Outer loop has  $O(V)$  iterations.

Implementing "**find**  $s \in vSet$  **with** minimum  $dist[s]$ "

1. try all  $s \in vSet \Rightarrow cost = O(V) \Rightarrow overall cost = O(E + V^2) = O(V^2)$
2. using a PQueue to implement extracting minimum
  - can improve overall cost to  $O(E + V \cdot \log V)$  (for best-known implementation)

## All-pair Shortest Path (APSP)

162/188


Given: weighted digraph  $G$

Result: shortest paths between all pairs of vertices

- $dist[][]$   $V \times V$ -indexed matrix of cost of shortest path from  $v_{row}$  to  $v_{col}$

- $path[][]$   $V \times V$ -indexed matrix of next node in shortest path from  $v_{row}$  to  $v_{col}$

Example:



V	0	1	2	3	4	
0	0	0.3	0.7	1.1	inf	dist
1	1.8	0	0.6	1.0	inf	
2	1.2	0.5	0	0.4	inf	
3	0.8	1.1	1.5	0	inf	
4	1.3	0.6	0.1	0.5	0	

V	0	1	2	3	4	
0	-	1	2	2	-	path
1	2	-	2	2	-	
2	3	1	-	3	-	
3	0	0	0	-	-	
4	2	2	2	2	-	

Shortest paths between all vertices

## Floyd's Algorithm

163/188

One approach to solving all-pair shortest path problem...

Data:  $G, dist[][], path[][]$  Algorithm:

$dist[][]$  // cost of shortest path from  $s$  to  $t$   
 $path[][]$  // next node after  $s$  on shortest path from  $s$  to  $t$

```
floydAPSP(G):
    Input graph G

    initialise dist[s][t]=0 for each s=t
                                   =w for each (s,t,w)∈edges(G)
                                   =∞ otherwise


    initialise path[s][t]=t for each (s,t,w)∈edges(G)
                                   =-1 otherwise

    for all i∈vertices(G) do
        for all s∈vertices(G) do
            for all t∈vertices(G) do
                if dist[s][i]+dist[i][t] < dist[s][t] then
                    dist[s][t]=dist[s][i]+dist[i][t]
                    path[s][t]=path[s][i]
                end if
            end for
        end for
    end for
```

### Exercise #24: Floyd's Algorithm

164/188

Show how Floyd's algorithm runs on:



### ... Floyd's Algorithm

Why Floyd's algorithm is correct:

A shortest path from  $s$  to  $t$  using only nodes from  $\{0, \dots, i\}$  is the shorter of

- a shortest path from  $s$  to  $t$  using only nodes from  $\{0, \dots, i-1\}$
- a shortest path from  $s$  to  $i$  using only nodes from  $\{0, \dots, i-1\}$  plus a shortest path from  $i$  to  $t$  using only nodes from  $\{0, \dots, i-1\}$



Also known as Floyd-Warshall algorithm (can you see why?)

### ... Floyd's Algorithm

Cost analysis ...

- initialising  $\text{dist}[\ ][\ ], \text{path}[\ ][\ ] \Rightarrow O(V^2)$
- $V$  iterations to update  $\text{dist}[\ ][\ ], \text{path}[\ ][\ ] \Rightarrow O(V^3)$


Time complexity of Floyd's algorithm:  $O(V^3)$  (same as Warshall's algorithm for transitive closure)

## Network Flow

### Exercise #25: Merchandise Distribution

Lucky Cricket Company ...

- produces cricket balls in Fairfield
- has a warehouse in Rozelle that stocks them
- ships them from factory to warehouse by leasing space on trucks with limited capacity:



What kind of algorithm would ...

- help us find the maximum number of crates that can be shipped from Fairfield to Rozelle per day?

## Flow Networks

Flow network ...

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	14	9	7			[0]		1	2	3		
[1]		0			5		[1]					4	
[2]		4	0			3	[2]		1				5
[3]			10	0		15	[3]			2			5
[4]					0		[4]						
[5]					2	0	[5]					4	

After 1<sup>st</sup> iteration  $i=0$ : unchanged After 2<sup>nd</sup> iteration  $i=1$ :

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	14	9	7	19	$\infty$	[0]	-	1	2	3	1	-
[1]	$\infty$	0	$\infty$	$\infty$	5	$\infty$	[1]	-	-	-	-	4	-
[2]	$\infty$	4	0	$\infty$	9	3	[2]	-	1	-	-	1	5
[3]	$\infty$	$\infty$	10	0	$\infty$	15	[3]	-	-	2	-	-	5
[4]	$\infty$	$\infty$	$\infty$	$\infty$	0	$\infty$	[4]	-	-	-	-	-	-
[5]	$\infty$	$\infty$	$\infty$	$\infty$	2	0	[5]	-	-	-	-	4	-

After 3<sup>rd</sup> iteration  $i=2$ :

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	13	9	7	18	12	[0]	-	2	2	3	2	2
[1]	$\infty$	0	$\infty$	$\infty$	5	$\infty$	[1]	-	-	-	-	4	-
[2]	$\infty$	4	0	$\infty$	9	3	[2]	-	1	-	-	1	5
[3]	$\infty$	14	10	0	19	13	[3]	-	2	2	-	2	2
[4]	$\infty$	$\infty$	$\infty$	$\infty$	0	$\infty$	[4]	-	-	-	-	-	-
[5]	$\infty$	$\infty$	$\infty$	$\infty$	2	0	[5]	-	-	-	-	4	-

After 4<sup>th</sup> iteration  $i=3$ : unchanged After 5<sup>th</sup> iteration  $i=4$ : unchanged After 6<sup>th</sup> iteration  $i=5$ :

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	13	9	7	14	12	[0]	-	2	2	3	2	2
[1]	$\infty$	0	$\infty$	$\infty$	5	$\infty$	[1]	-	-	-	-	4	-
[2]	$\infty$	4	0	$\infty$	5	3	[2]	-	1	-	-	5	5
[3]	$\infty$	14	10	0	15	13	[3]	-	2	2	-	2	2
[4]	$\infty$	$\infty$	$\infty$	$\infty$	0	$\infty$	[4]	-	-	-	-	-	-
[5]	$\infty$	$\infty$	$\infty$	$\infty$	2	0	[5]	-	-	-	-	4	-



- weighted graph  $G=(V,E)$
- distinct nodes  $s \in V$  (source),  $t \in V$  (sink)

Edge weights denote *capacities*

Applications:

- Distribution networks, e.g.
  - source: oil field
  - sink: refinery
  - edges: pipes
- Traffic flow

## ... Flow Networks

171/188

Flow in a network  $G=(V,E)$  ... nonnegative  $f(v,w)$  for all vertices  $v,w \in V$  such that

- $f(v,w) \leq \text{capacity}$  for each edge  $e=(v,w, \text{capacity}) \in E$
- $f(v,w)=0$  if no edge between  $v$  and  $w$
- total flow *into* a vertex = total flow *out of* a vertex:


$$\sum_{x \in V} f(x, v) = \sum_{y \in V} f(v, y) \quad \text{for all } v \in V \setminus \{s, t\}$$

Maximum flow ... no other flow from  $s$  to  $t$  has larger value


## ... Flow Networks

172/188

Example:



A (maximum) flow ...




## Augmenting Paths

173/188

Assume ...  $f(v,w)$  contains current flow

*Augmenting path*: any path from source  $s$  to sink  $t$  that can currently take more flow

Example:



## Residual Network


174/188

Assume ... flow network  $G=(V,E)$  and flow  $f(v,w)$

*Residual network*  $(V,E')$ :

- same vertex set  $V$
- for each edge  $v \xrightarrow{c} w \in E$  ...
  - $f(v,w) < c \Rightarrow$  add edge  $(v \xrightarrow{c-f(v,w)} w)$  to  $E'$
  - $f(v,w) > 0 \Rightarrow$  add edge  $(v \xleftarrow{f(v,w)} w)$  to  $E'$


Example:



## Exercise #26: Augmenting Paths and Residual Networks


175/188

Find an augmenting path in:




and show the residual network after augmenting the flow

### 1. Augmenting path:



maximum additional flow = 1

2. Residual network:



Can you find a further augmenting path in the new residual network?

Edmonds-Karp Algorithm

177/188

One approach to solving maximum flow problem ...

maxflow(G):

- 1. Find a shortest augmenting path
- 2. Update flow[ ][ ] so as to represent residual network
- 3. Repeat until no augmenting path can be found

... Edmonds-Karp Algorithm

178/188

Algorithm:

flow[ ][ ] // VxV array of current flow  
visited[ ] /\* array of predecessor nodes on shortest path  
            from source to sink in residual network \*/

```
maxflow(G):
  Input  flow network G with source s and sink t
  Output maximum flow value

  initialise flow[v][w]=0 for all vertices v, w
  maxflow=0

  while 3shortest augmenting path from s to t do /* Run BFS on "residual network"
    given by capacity[v][w] > flow[v][w]
    to find a shortest path "visited[]" */
    df = maximum additional flow via visited[]
    // adjust flow so as to represent residual graph
    v=t
    while v≠s do
      flow[visited[v]][v] = flow[visited[v]][v] + df;
      flow[v][visited[v]] = flow[v][visited[v]] - df;
      v=visited[v]
    end while
    maxflow=maxflow+df
  end while
  return maxflow
```

... Edmonds-Karp Algorithm

179/188

Time complexity analysis ...

- Theorem. The number of augmenting paths needed is at most  $V \cdot E/2$ .  
⇒ Outer loop has  $O(V \cdot E)$  iterations.
- Finding augmenting path ⇒  $O(E)$  (consider only vertices connected to source and sink ⇒  $O(V+E)=O(E)$ )


Overall cost of Edmonds-Karp algorithm:  $O(V \cdot E^2)$

Note: Edmonds-Karp algorithm is an implementation of general Ford-Fulkerson method

Exercise #27: Edmonds-Karp Algorithm

180/188

Show how Edmonds-Karp algorithm runs on:



flow	[0]	[1]	[2]	[3]	[4]	[5]	c>f?	[0]	[1]	[2]	[3]	[4]	[5]
[0]							[0]						
[1]							[1]						
[2]							[2]						
[3]							[3]						
[4]							[4]						
[5]							[5]						

flow	[0]	[1]	[2]	[3]	[4]	[5]	c>f?	[0]	[1]	[2]	[3]	[4]	[5]	df	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	0	0	0	0	0	[0]		✓	✓				[0]		2	3			
[1]	0	0	0	0	0	0	[1]				✓	✓		[1]				3	1	
[2]	0	0	0	0	0	0	[2]				✓	✓		[2]				1	1	
[3]	0	0	0	0	0	0	[3]						✓	[3]						2
[4]	0	0	0	0	0	0	[4]						✓	[4]						3
[5]	0	0	0	0	0	0	[5]							[5]						

augmenting path: 0-1-3-5, df: 2

flow	[0]	[1]	[2]	[3]	[4]	[5]	c>f?	[0]	[1]	[2]	[3]	[4]	[5]	df	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	0	0	0	0	[0]			✓				[0]			3			
[1]	-2	0	0	2	0	0	[1]	✓			✓	✓		[1]	2			1	1	
[2]	0	0	0	0	0	0	[2]				✓	✓		[2]				1	1	
[3]	0	-2	0	0	0	2	[3]		✓					[3]		2				
[4]	0	0	0	0	0	0	[4]						✓	[4]						3
[5]	0	0	0	-2	0	0	[5]				✓			[5]				2		

augmenting path: 0-2-4-5, df: 1

f <sub>low</sub>	[0]	[1]	[2]	[3]	[4]	[5]	c>f?	[0]	[1]	[2]	[3]	[4]	[5]	df	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	1	0	0	0	[0]			✓				[0]			2			
[1]	-2	0	0	2	0	0	[1]	✓			✓	✓		[1]	2			1	1	
[2]	-1	0	0	0	1	0	[2]	✓			✓			[2]	1			1		
[3]	0	-2	0	0	0	2	[3]		✓					[3]		2				
[4]	0	0	-1	0	0	1	[4]			✓			✓	[4]			1			2
[5]	0	0	0	-2	-1	0	[5]				✓	✓		[5]				2	1	

augmenting path: 0-2-3-1-4-5, df: 1

f <sub>low</sub>	[0]	[1]	[2]	[3]	[4]	[5]	c>f?	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	2	0	0	0	[0]			✓			
[1]	-2	0	0	1	1	0	[1]	✓			✓		
[2]	-2	0	0	1	1	0	[2]	✓					
[3]	0	-1	-1	0	0	2	[3]		✓	✓			
[4]	0	-1	-1	0	0	2	[4]		✓	✓			✓
[5]	0	0	0	-2	-2	0	[5]				✓	✓	

## Digraph Applications

### PageRank

183/188

**Goal:** determine which Web pages are "important"

**Approach:** ignore page contents; focus on hyperlinks

- treat Web as graph: page = vertex, hyperlink = directed edge
- pages with many incoming hyperlinks are important
- need to compute "incoming degree" for vertices

Problem: the Web is a *very* large graph

- approx.  $10^{14}$  pages,  $10^{15}$  hyperlinks

Assume for the moment that we could build a graph ...

Most frequent operation in algorithm "Does edge (v,w) exist?"

#### ... PageRank

184/188

Simple PageRank algorithm:

```
PageRank(myPage):
  rank=0
  for each page in the Web do
    if linkExists(page,myPage) then
      rank=rank+1
    end if
  end for
```

Note: requires *inbound* link check

#### ... PageRank

185/188

$V$  = # pages in Web,  $E$  = # hyperlinks in Web

Costs for computing PageRank for each representation:

Representation	linkExists(v,w)	Cost
Adjacency <i>matrix</i>	edge[v][w]	$I$
Adjacency <i>lists</i>	inLL(list[v],w)	$\approx E/V$

Not feasible ...

- adjacency matrix ...  $V \approx 10^{14} \Rightarrow$  matrix has  $10^{28}$  cells
- adjacency list ...  $V$  lists, each with  $\approx 10$  hyperlinks  $\Rightarrow 10^{15}$  list nodes

So how to really do it?

#### ... PageRank

186/188

Approach: the random web surfer

- if we randomly follow links in the web ...
- ... more likely to re-discover pages with many inbound links

```
curr=random page, prev=null
for a long time do
  if curr not in array ranked[] then
    rank[curr]=0
  end if
  rank[curr]=rank[curr]+1
  if random(0,100)<85 then // with 85% chance ...
    prev=curr
    curr=choose hyperlink from curr // ... crawl on
  else
    curr=random page // avoid getting stuck
    prev=null
  end if
end for
```

Could be accomplished while we crawl web to build search index

187/188

## Exercise #28: Implementing Facebook

Facebook could be considered as a giant "social graph"

- what are the vertices?
- what are the edges?
- are edges directional?

What kind of algorithm would ...

- help us find people that you might like to "befriend"?

---

## Summary

188/188

- Graph terminology
  - vertices, edges, vertex degree, connected graph, tree
  - path, cycle, clique, spanning tree, spanning forest
- Graph representations
  - array of edges
  - adjacency matrix
  - adjacency lists
- Graph traversal
  - depth-first search (DFS)
  - breadth-first search (BFS)
  - cycle check, connected components
  - Hamiltonian paths/circuits, Euler paths/circuits
- Digraphs, weighted graphs: representations, applications
- Reachability
  - Warshall
- Minimum Spanning Tree (MST)
  - Kruskal, Prim
- Shortest path problems
  - Dijkstra (single source SPP)
  - Floyd (all-pair SPP)
- Flow networks
  - Edmonds-Karp (maximum flow)
  
- Suggested reading (Sedgewick):
  - graph representations ... Ch. 17.1-17.5
  - Hamiltonian/Euler paths ... Ch. 17.7
  - graph search ... Ch. 18.1-18.3, 18.7
  - digraphs ... Ch. 19.1-19.3
  - weighted graphs ... Ch. 20-20.1
  - MST ... Ch. 20.2-20.4
  - SSP ... Ch. 21-21.3
  - network flows ... Ch. 22.1-22.2