

Searching

4/129

An extremely common application in computing

- given a (large) collection of *items* and a *key* value
- find the item(s) in the collection containing that key
 - item = (key, val₁, val₂, ...) (i.e. a structured data type)
 - key = value used to distinguish items (e.g. student ID)

Applications: Google, databases,

... Searching

5/129

Since searching is a very important/frequent operation, many approaches have been developed to do it

- Linear structures: arrays, linked lists
- Arrays = random access. Lists = sequential access

Cost of searching:

	Array	List
Unsorted	O(n) (linear scan)	O(n) (linear scan)
Sorted	O(log n) (binary search)	O(n) (linear scan)

- $O(n)$... linear scan (search technique of last resort)
- $O(\log n)$... binary search, *search trees* (trees also have other uses)

Also (cf. Sedgewick Ch.14): hash tables ($O(1)$), but only under optimal conditions)


... Searching

6/129

Maintaining the order in sorted arrays and files is a costly operation.

Search trees are as efficient to search but more efficient to maintain.

Example: the following tree corresponds to the sorted array [2, 5, 10, 12, 14, 17, 20, 24, 29, 30, 31, 32]:




Tree Data Structures

7/129

Trees are connected graphs

- consisting of nodes and edges (called *links*), with no cycles (no "up-links")
- each node contains a **data** value (or key+data)
- each node has **links** to $\leq k$ other child nodes ($k=2$ below)




... Tree Data Structures

8/129

Trees can be used as a data structure, but also for *illustration*.

E.g. showing evaluation of a prefix arithmetic expression




... Tree Data Structures

9/129

Binary trees ($k=2$ children per node) can be defined recursively, as follows:

A *binary tree* is either

- empty (contains no nodes)
- consists of a *node*, with two *subtrees*
 - node contains a value
 - left and right subtrees are *binary trees*




... Tree Data Structures

10/129

Other special kinds of tree

- *m-ary tree*: each internal node has exactly m children
- *Ordered tree*: all left values $<$ root, all right values $>$ root
- *Balanced tree*: has \approx minimal height for a given number of nodes
- *Degenerate tree*: has \approx maximal height for a given number of nodes



Perfectly balanced binary trees have the properties

- #nodes in left subtree = #nodes in right subtree
- this property applies over all nodes in the tree

Shape of tree is determined by order of insertion.

Search Trees


11/129

Binary Search Trees

12/129

Binary search trees (or *BSTs*) have the characteristic properties

- each node is the root of 0, 1 or 2 subtrees
- all values in any left subtree are less than root
- all values in any right subtree are greater than root
- these properties applies over all nodes in the tree



... Binary Search Trees

13/129

Operations on BSTs:

- *insert*(Tree,Item) ... add new item to tree via key
- *delete*(Tree,Key) ... remove item with specified key from tree
- *search*(Tree,Key) ... find item containing key in tree
- plus, "bookkeeping" ... new(), free(), show(), ...

Notes:


- in general, nodes contain *Items*; we just show *Item.key*
- keys are unique (not technically necessary)

... Binary Search Trees

14/129

Level of node = path length from root to node

Height (or: *depth*) of tree = max path length from root to leaf



Height-balanced tree: \forall nodes: $\text{height}(\text{left subtree}) = \text{height}(\text{right subtree}) \pm 1$

Time complexity of tree algorithms is typically $O(\text{height})$

Exercise #1: Insertion into BSTs

15/129

For each of the sequences below

- start from an initially empty binary search tree

- show tree resulting from inserting values in order given

(a) 4 2 6 5 1 7 3

(b) 6 5 2 3 4 7 1

(c) 1 2 3 4 5 6 7

Assume new values are always inserted as new leaf nodes

(a) the balanced tree on slide 10 (height = 2)

(b) the non-balanced tree on slide 10 (height = 4)

(c) a fully degenerate tree of height 6

Representing BSTs


17/129

Binary trees are typically represented by node structures

- containing a value, and pointers to child nodes


Most tree algorithms move *down* the tree.

If upward movement needed, add a pointer to parent.



... Representing BSTs

18/129



Typical data structures for trees ...

```
// a Tree is represented by a pointer to its root node
typedef struct Node *Tree;
```

```
// a Node contains its data, plus left and right subtrees
typedef struct Node {
    Item data;           // We will only use an int for the value of a node
    Tree left, right;
} Node;
```

```
// some macros that we will use frequently
#define data(tree) ((tree)->data)
#define left(tree) ((tree)->left)
#define right(tree) ((tree)->right)
```

We ignore items \Rightarrow data in Node is just a key

Tree Algorithms

Searching in BSTs

20/129

Most tree algorithms are best described recursively

TreeSearch(tree,item):

Input tree, item
Output true if item found in tree, false otherwise

```
if tree is empty then
    return false
else if item < data(tree) then
    return TreeSearch(left(tree),item)
else if item > data(tree) then
    return TreeSearch(right(tree),item)
else // found
    return true
end if
```

Insertion into BSTs

21/129

Insert an item into appropriate subtree

insertAtLeaf(tree,item):

Input tree, item
Output tree with item inserted

```
if tree is empty then
    return new node containing item
else if item < data(tree) then
    return insertAtLeaf(left(tree),item)
else if item > data(tree) then
    return insertAtLeaf(right(tree),item)
else
    return tree // avoid duplicates
end if
```

Tree Traversal

22/129

Iteration (traversal) on ...

- **Lists** ... visit each value, from first to last
- **Graphs** ... visit each vertex, order determined by DFS/BFS/...


For binary **Trees**, several well-defined visiting orders exist:

- *preorder* (NLR) ... visit root, then left subtree, then right subtree
- *inorder* (LNR) ... visit left subtree, then root, then right subtree
- *postorder* (LRN) ... visit left subtree, then right subtree, then root
- *level-order* ... visit root, then all its children, then all their children

... Tree Traversal

23/129

Consider "visiting" an expression tree like:




NLR: + * 1 3 - * 5 7 9 (prefix-order: useful for building tree)
LNR: 1 * 3 + 5 * 7 - 9 (infix-order: "natural" order)
LRN: 1 3 * 5 7 * 9 - + (postfix-order: useful for evaluation)
Level: + * - 1 3 * 9 5 7 (level-order: useful for printing tree)

Exercise #2: Tree Traversal

24/129

Show NLR, LNR, LRN traversals for the tree



NLR (preorder): 20 10 5 2 14 12 17 30 24 29 32 31

LNR (inorder): 2 5 10 12 14 17 20 24 29 30 31 32

LRN (postorder): 2 5 12 17 14 10 29 24 31 32 30 20

Exercise #3: Non-recursive traversals

26/129

Write a non-recursive *preorder* traversal algorithm.

Assume that you have a stack ADT available.

```
showBSTreePreorder(t):  
    Input tree t  
  
    push t onto new stack S  
    while stack is not empty do  
        t=pop(S)  
        print data(t)  
        if right(t) is not empty then  
            push right(t) onto S  
        end if  
        if left(t) is not empty then  
            push left(t) onto S  
        end if  
    end while
```

Joining Two Trees

28/129

An auxiliary tree operation ...


Tree operations so far have involved just one tree.

An operation on two trees: $t = \text{joinTrees}(t_1, t_2)$

- Pre-condition:
 - $\max(\text{key}(t_1)) < \min(\text{key}(t_2))$
- Post-condition:
 - result is a BST (i.e. correctly ordered) with all items from t_1 and t_2

Method:

- find the min node in the right subtree (t_2)
- replace min node by its right subtree
- elevate min node to be new root of both trees



Advantage: doesn't increase height of tree significantly

$x \leq \text{height}(t) \leq x+1$, where $x = \max(\text{height}(t_1), \text{height}(t_2))$

... Joining Two Trees

29/129

Implementation of tree-join

```

joinTrees( $t_1, t_2$ ):
  Input  trees  $t_1, t_2$ 
  Output  $t_1$  and  $t_2$  joined together



  if  $t_1$  is empty then return  $t_2$ 
  else if  $t_2$  is empty then return  $t_1$ 
  else
    curr= $t_2$ , parent=NULL
    while left(curr) is not empty do      // find min element in  $t_2$ 
      parent=curr
      curr=left(curr)
    end while
    if parent≠NULL then
      left(parent)=right(curr) // unlink min element from parent
      right(curr)= $t_2$ 
    end if
    left(curr)= $t_1$ 
    return curr                // curr is new root
  end if

```

Exercise #4: Joining Two Trees

30/129

Join the trees



Deletion from BSTs

32/129

Insertion into a binary search tree is easy.

Deletion from a binary search tree is harder.


Four cases to consider ...

- empty tree ... new tree is also empty
- zero subtrees ... unlink node from parent
- one subtree ... replace by child
- two subtrees ... replace by successor, join two subtrees


... Deletion from BSTs

33/129

Case 2: item to be deleted is a leaf (zero subtrees)




Just delete the item:




... Deletion from BSTs

34/129

Case 3: item to be deleted has one subtree




Replace the item by its only subtree:




... Deletion from BSTs

35/129

Case 4: item to be deleted has two subtrees




Version 1: right child becomes new root, attach left subtree to min element of right subtree




... Deletion from BSTs

36/129

Case 4: item to be deleted has two subtrees



Version 2: *join* left and right subtree



Advantage: doesn't increase height of tree significantly

... Deletion from BSTs

37/129

Pseudocode (with version 2 for case 4)

```

TreeDelete(t,item):
    Input  tree t, item
    Output t with item deleted

    if t is not empty then           // nothing to do for case 1
        if item < data(t) then      // delete item in left subtree
            left(t)=TreeDelete(left(t),item)
        else if item > data(t) then  // delete item in right subtree
            right(t)=TreeDelete(right(t),item)
        else                         // node 't' must be deleted
            if left(t) and right(t) are empty then
                new=empty tree      // case 2: 0 children
            else if left(t) is empty then
                new=right(t)        // case 3: 1 child
            else if right(t) is empty then
                new=left(t)         // case 3: 1 child
            else
                new=joinTrees(left(t),right(t)) // case 4: 2 children
            end if
            free memory allocated for current node t
            t=new
        end if
    end if
    return t

```

Balanced Search Trees

Balanced BSTs


39/129

Goal: build binary search trees which have

- minimum height \Rightarrow minimum worst case search cost

Best balance you can achieve for tree with N nodes:

- $\text{abs}(\# \text{nodes}(\text{LeftSubtree}) - \# \text{nodes}(\text{RightSubtree})) \leq 1$, for every node
- height of $\log_2 N \Rightarrow$ worst case search $O(\log N)$



... Balanced BSTs

40/129

To assist with rebalancing, we consider new operations:

Left rotation

- move right child to root; rearrange links to retain order

Right rotation

- move left child to root; rearrange links to retain order


Insertion at root

- each new item is added as the new root node

Operation for Rebalancing: Tree Rotation

41/129

In tree below: $t_1 < n_2 < t_2 < n_1 < t_3$



Method for rotating tree T right:

- N_1 is current root; N_2 is root of N_1 's left subtree
- N_1 gets new left subtree, which is N_2 's right subtree
- N_1 becomes root of N_2 's new right subtree
- N_2 becomes new root

Left rotation: swap left/right in the above.

Cost of tree rotation: $O(1)$

... Operation for Rebalancing: Tree Rotation


42/129

Algorithm for right rotation:

```

rotateRight( $n_1$ ):
|   Input   tree  $n_1$ 
|   Output   $n_1$  rotated to the right
|
|   if  $n_1$  is empty or left( $n_1$ ) is empty then
|       return  $n_1$ 
|   end if
|    $n_2 = \text{left}(n_1)$ 
|   left( $n_1$ ) = right( $n_2$ )
|   right( $n_2$ ) =  $n_1$ 
|   return  $n_2$ 


```




Exercise #5: Tree Rotation

43/129

Consider the tree t :




Show the result of rotateRight(t)



Exercise #6: Tree Rotation

45/129

Write the algorithm for left rotation



```
rotateLeft(n2):
  Input tree n2
  Output n2 rotated to the left

  if n2 is empty or right(n2) is empty then
    return n2
  end if
  n1=right(n2)
  right(n2)=left(n1)
  left(n1)=n2
  return n1
```

Insertion at Root

47/129

Previous description of BSTs inserted at leaves.

Different approach: insert new item at root.

Potential disadvantages:

- large-scale rearrangement of tree for each insert

Potential advantages:

- recently-inserted items are close to root
- low cost if recent items more likely to be searched

... Insertion at Root


48/129

Method for inserting at root:

- base case:
 - tree is empty; make new node and make it root
- recursive case:
 - insert new node as root of appropriate subtree
 - lift new node to root by rotation

... Insertion at Root


49/129




Exercise #7: Insertion at Root

50/129

Consider the tree t:



Show the result of insertAtRoot(t, 24)



... Insertion at Root

52/129

Analysis of insertion-at-root:

- same complexity as for insertion-at-leaf: $O(height)$
- tendency to be balanced, but no balance guarantee
- benefit comes in searching
 - for some applications, search favours recently-added items
 - insertion-at-root ensures these are close to root
- could even consider "move to root when found"
 - effectively provides "self-tuning" search tree

⇒ more on this later (real balanced trees)

Rebalancing Trees

Tree Review

60/129

Binary search trees ...


- data structures designed for $O(\log n)$ search
- consist of nodes containing item (incl. key) and two links
- can be viewed as recursive data structure (subtrees)
- have overall ordering (data(Left) < root < data(Right))
- insert new nodes as leaves (or as root), delete from anywhere
- have structure determined by insertion order (*worst*: $O(n)$)
- operations: insert, delete, search, ...

Balanced BSTs

61/129

Reminder ...

- Goal: build binary search trees which have
 - minimum height \Rightarrow minimum worst case search cost
- Best balance you can achieve for tree with N nodes:
 - tree height of $\log_2 N \Rightarrow$ worst case search $O(\log N)$



Randomised BST Insertion

62/129

Effects of order of insertion on BST shape:

- best case (for at-leaf insertion): keys inserted in pre-order (median key first, then median of lower half, median of upper half, etc.)
- worst case: keys inserted in ascending/descending order
- average case: keys inserted in *random* order $\Rightarrow O(\log_2 n)$

Tree ADT has no control over order that keys are supplied.

Can the algorithm itself introduce some *randomness*?

In the hope that this randomness helps to balance the tree ...

... Randomised BST Insertion

63/129

How can a computer pick a number at random?

- it cannot

Software can only produce *pseudo random numbers*.

- a pseudo random number may appear unpredictable
 - but is actually predictable
- \Rightarrow implementation may deviate from expected theoretical behaviour
 - more on this in week 5

... Randomised BST Insertion

64/129

- Pseudo random numbers in C:

```
rand() // generates random numbers in the range 0 .. RAND_MAX
```

where the constant `RAND_MAX` is defined in `stdlib.h`
(depends on the computer: on the CSE network, `RAND_MAX` = 2147483647)

To convert the return value of `rand()` to a number between 0 .. RANGE

- compute the remainder after division by `RANGE+1`

... Randomised BST Insertion

65/129

Approach: normally do leaf insert, randomly do root insert.

```
insertRandom(tree,item)
| Input tree, item
| Output tree with item randomly inserted
|
| if tree is empty then
|   return new node containing item
| end if
| // p/q chance of doing root insert
| if random number mod q < p then
|   return insertAtRoot(tree,item)
| else
|   return insertAtLeaf(tree,item)
| end if
```

E.g. 30% chance \Rightarrow choose $p=3, q=10$

... Randomised BST Insertion

66/129

Cost analysis:

- similar to cost for inserting keys in random order: $O(\log n)$
- does not rely on keys being supplied in random order

Approach can also be applied to deletion:

- standard method promotes inorder successor to root
- for the randomised method ...
 - promote inorder successor from right subtree, OR
 - promote inorder predecessor from left subtree

Rebalancing Trees

67/129

Another approach to balanced trees:

- insert into leaves as for simple BST
- periodically, rebalance the tree

Question: how frequently/when/how to rebalance?

```
NewTreeInsert(tree,item):
  Input  tree, item
  Output tree with item randomly inserted

  t=insertAtLeaf(tree,item)
  if #nodes(t) mod k = 0 then
    t=rebalance(t)
  end if
  return t
```

E.g. rebalance after every 20 insertions \Rightarrow choose $k=20$


Note: To do this efficiently we would need to change tree data structure and basic operations:

```
typedef struct Node {
  Item data;
  int  nnodes; // #nodes in my tree
  Tree left, right; // subtrees
} Node;
```

... Rebalancing Trees

68/129

How to rebalance a BST? Move median item to root.



Implementation of rebalance:

```
rebalance(t):
  Input  tree t with n nodes
  Output t rebalanced


  if n≥3 then
    t=partition(t,⌊n/2⌋) // put node with median key at root
    left(t)=rebalance(left(t)) // then rebalance each subtree
    right(t)=rebalance(right(t))
  end if
  return t
```

... Rebalancing Trees

69/129

New operation on trees:

- **partition(tree,i)**: re-arrange tree so that element with index i becomes root




For tree with n nodes, indices are $0 \dots n-1$

Exercise #8: Partition

70/129

Consider this tree with $n = 6$ nodes:




Which element has index $\lfloor n/2 \rfloor = 3$?

29

... Rebalancing Trees

72/129

Partition: moves i^{th} node to root



Algorithm:

```

partition(tree, i):
  Input  tree with n nodes, index i
  Output tree with item #i moved to the root

  m=#nodes(left(tree))
  if i < m then
    left(tree)=partition(left(tree), i)
    tree=rotateRight(tree)
  else if i > m then
    right(tree)=partition(right(tree), i-m-1)
    tree=rotateLeft(tree)
  end if
  return tree


```

Note: $\text{size}(\text{tree}) = n$, $\text{size}(\text{left}(\text{tree})) = m$, $\text{size}(\text{right}(\text{tree})) = n-m-1$ (why -1?)


Exercise #9: Partition

73/129

Consider the tree t :



Show the result of $\text{partition}(t, 3)$



... Rebalancing Trees

Even the most efficient implementation of rebalancing requires (in the worst case) to visit every node $\Rightarrow O(N)$

Cost means not feasible to rebalance after each insertion.

When to rebalance? ... Some possibilities:

- after every k insertions
- whenever "imbalance" exceeds threshold

Either way, we tolerate worse search performance for periods of time.

Does it solve the problem? ... Not completely \Rightarrow Solution: real balanced trees (later)

Splay Trees

77/129

A kind of "self-balancing" tree ...

Splay tree insertion modifies insertion-at-root method:

- by considering *parent-child-grandchild* (three level analysis)
- by performing double-rotations based on p-c-g orientation

The idea: appropriate double-rotations improve tree balance.

Splay tree implementations also do *rotation-in-search*:

- by performing double-rotations also when searching

The idea: provides similar effect to periodic rebalance.


\Rightarrow improves balance but makes search more expensive

... Splay Trees


78/129

Cases for splay tree double-rotations:

- case 1: grandchild is left-child of left-child \Rightarrow double right rotation from top
- case 2: grandchild is right-child of left-child
- case 3: grandchild is left-child of right-child
- case 4: grandchild is right-child of right-child \Rightarrow double left rotation from top




Double-rotation case for left-child of left-child ("zig-zig"):




Similarly for right-child of right-child ("zag-zag")

Note: both rotations at the root (unlike insertion-at-root)


Example:




Tree after "zag-zag" rotation:



vs. promoting 36 to the root (a la insertion-at-root):



Double-rotation case for right-child of left-child ("zig-zag"):



Similarly for left-child of right-child ("zag-zig")

Note: rotate subtree first (like insertion-at-root)

Algorithm for splay tree insertion:

```
insertSplay(tree,item):
  Input  tree, item
  Output tree with item splay-inserted



  if tree is empty then return new node containing item
  else if item=data(tree) then return tree
  else if item<data(tree) then
    if left(tree) is empty then
      left(tree)=new node containing item
    else if item<data(left(tree)) then
      // Case 1: left-child of left-child "zig-zig"
      left(left(tree))=insertSplay(left(left(tree)),item)
      tree=rotateRight(tree)
    else if item>data(left(tree)) then
      // Case 2: right-child of left-child "zig-zag"
      right(left(tree))=insertSplay(right(left(tree)),item)
      left(tree)=rotateLeft(left(tree))
    end if
    return rotateRight(tree)
  else // item>data(tree)
    if right(tree) is empty then
      right(tree)=new node containing item
    else if item<data(right(tree)) then
      // Case 3: left-child of right-child "zag-zig"
      left(right(tree))=insertSplay(left(right(tree)),item)
      right(tree)=rotateRight(right(tree))
    else if item>data(right(tree)) then
      // Case 4: right-child of right-child "zag-zag"
      right(right(tree))=insertSplay(right(right(tree)),item)
      tree=rotateLeft(tree)
    end if
    return rotateLeft(tree)
```

| end if

Exercise #10: Splay Trees

83/129

Insert 36 into this splay tree:



... Splay Trees

85/129

Searching in splay trees:

```
searchSplay(tree,item):
  Input  tree, item
  Output address of item if found in tree
         NULL otherwise


  if tree=NULL then
    return NULL
  else
    tree=splay(tree,item)
    if data(tree)=item then
      return tree
    else
      return NULL
    end if
  end if
```

where `splay()` is similar to `insertSplay()`, except that it doesn't add a node ... simply moves `item` to root if found, or nearest node if not found


... Splay Trees

86/129

Example:



Splay tree after searching for 22:



... Splay Trees

87/129

Why take into account both child and grandchild?

- moves accessed node to the root
- *moves every ancestor of accessed node roughly halfway to the root*

⇒ better amortized cost than insert-at-root

Analysis of splay tree performance:

- assume that we "splay" for both insert and search
- consider: m insert+search operations, n nodes
- *Theorem.* Total number of comparisons: average $O((n+m) \cdot \log(n+m))$

Gives good overall (amortized) cost.

- insert cost not significantly different to insert-at-root
- search cost increases, but ...
 - improves balance on each search
 - moves frequently accessed nodes closer to root

But ... still has worst-case search cost $O(n)$

Real Balanced Trees

Better Balanced Binary Search Trees

89/129

So far, we have seen ...

- occasional rebalance ... fix balance periodically
- splay trees ... reasonable amortized performance
- but both types still have $O(n)$ worst case

Ideally, we want both average/worst case to be $O(\log n)$

- AVL trees ... fix imbalances as soon as they occur
- 2-3-4 trees ... use varying-sized nodes to assist balance
- red-black trees ... isomorphic to 2-3-4, but binary nodes

AVL Trees

90/129

Invented by Georgy Adelson-Velsky and Evgenii Landis

Approach:

- insertion (at leaves) may cause imbalance
- repair balance as soon as we notice imbalance
- repairs done locally, not by overall tree restructure

A tree is unbalanced when: $\text{abs}(\text{height}(\text{left}) - \text{height}(\text{right})) > 1$

This can be repaired by at most two rotations:

- if left subtree too deep ...
 - if data inserted in left-right grandchild \Rightarrow left-rotate left subtree
 - rotate right
- if right subtree too deep ...
 - if data inserted in right-left grandchild \Rightarrow right-rotate right subtree
 - rotate left

Problem: determining height/depth of subtrees may be expensive.

... AVL Trees

91/129

Implementation of AVL insertion

```
insertAVL(tree,item):
  Input  tree, item
  Output tree with item AVL-inserted



  if tree is empty then
    return new node containing item
  else if item=data(tree) then
    return tree
  else
    if item<data(tree) then
      left(tree)=insertAVL(left(tree),item)
    else if item>data(tree) then
      right(tree)=insertAVL(right(tree),item)
    end if
    if height(left(tree))-height(right(tree)) > 1 then
      if item>data(left(tree)) then
        left(tree)=rotateLeft(left(tree))
      end if
      tree=rotateRight(tree)
    else if height(right(tree))-height(left(tree)) > 1 then
      if item<data(right(tree)) then
        right(tree)=rotateRight(right(tree))
      end if
      tree=rotateLeft(tree)
```

```
    end if
    return tree
  end if
```

Exercise #11: AVL Trees

92/129

Insert 27 into the AVL tree



What would happen if you now insert 28?

You may like the animation at www.cs.usfca.edu/~galles/visualization/AVLtree.html

... AVL Trees

94/129

Analysis of AVL trees:

- trees are *height*-balanced; subtree depths differ by +/-1
- average/worst-case search performance of $O(\log n)$
- *require* extra data to be stored in each node ("height")
- may not be *weight*-balanced; subtree sizes may differ




2-3-4 Trees

95/129

2-3-4 trees have three kinds of nodes

- 2-nodes, with two children (same as normal BSTs)
- 3-nodes, two values and three children
- 4-nodes, three values and four children



... 2-3-4 Trees

96/129

2-3-4 trees are ordered similarly to BSTs



In a *balanced* 2-3-4 tree:

- all leaves are at same distance from the root

... 2-3-4 Trees

97/129

Possible 2-3-4 tree data structure:

```
typedef struct node {
    int      degree;    // 2, 3 or 4
    int      data[3];   // items in node
    struct node *child[4]; // links to subtrees
} node;
```

... 2-3-4 Trees

98/129

Searching in 2-3-4 trees:

Search(tree, item):

```
Input   tree, item
Output address of item if found in 2-3-4 tree
         NULL otherwise

if tree is empty then
    return NULL
else
    i=0
    while i<tree.degree-1 and item>tree.data[i] do
        i=i+1 // find relevant slot in data[]
    end while
    if item=tree.data[i] then // data[i] exists and equals item
        return address of tree.data[i] // => item found
    else // keep looking in relevant subtree
        return Search(tree.child[i], item)
    end if
end if
```

... 2-3-4 Trees


2-3-4 tree searching cost analysis:

- as for other trees, worst case determined by height h
- 2-3-4 trees are always balanced \Rightarrow height is $O(\log n)$
- worst case for height: all nodes are 2-nodes
same case as for balanced BSTs, i.e. $h \approx \log_2 n$
- best case for height: all nodes are 4-nodes
balanced tree with branching factor 4, i.e. $h \approx \log_4 n$

... 2-3-4 Trees

100/129

Insertion into a 2-node or 3-node:




Insertion into a 4-node (requires a split):



... 2-3-4 Trees

101/129

2-3-4 trees grow "upwards" by splitting the root:




Exercise #12: 2-3-4 Trees

102/129

Insert C into this 2-3-4 tree:





... 2-3-4 Trees

104/129

Starting with the root node:

repeat


- if current node is full (i.e. contains 3 items)
 - split into two 2-nodes
 - promote middle element to parent
 - if no parent \Rightarrow middle element becomes the new root 2-node
 - go back to parent node
- if current node is a leaf
 - insert Item in this node, degree++
- if current node is not a leaf
 - go to child where Item belongs

until Item inserted

... 2-3-4 Trees

105/129



Building a 2-3-4 tree ... 7 insertions:



Exercise #13: 2-3-4 Tree Insertions

106/129

Show what happens when D, S, F, U are inserted into this tree:



... 2-3-4 Trees

108/129

Insertion algorithm:

```
insert(tree,item):
  Input  2-3-4 tree, item
  Output tree with item inserted

  node=root(tree), parent=NULL
  repeat
    if node.degree=4 then
      promote = node.data[1] // middle value
      nodeL   = new node containing node.data[0]
      nodeR   = new node containing node.data[2]
      if parent=NULL then
        make new 2-node root with promote,nodeL,nodeR
      else
        insert promote,nodeL,nodeR into parent
        increment parent.degree
      end if
      node=parent
    end if
    if node is a leaf then
      insert item into node
      increment node.degree
    else
      parent=node
      i=0
      while i<node.degree-1 and item>node.data[i] do
        i=i+1 // find relevant child to insert item
      end while
      node=node.child[i]
    end if
  until item inserted
```

... 2-3-4 Trees

109/129

Variations on 2-3-4 trees ...

Variation #1: why stop at 4? why not 2-3-4-5 trees? or M -way trees?

- allow nodes to hold up to $M-1$ items, and at least $M/2$
- if each node is a disk-page, then we have a B -tree (databases)
- for B-trees, depending on Item size, $M > 100/200/400$

Variation #2: don't have "variable-sized" nodes

- use standard BST nodes, augmented with one extra piece of data
- implement similar strategy as 2-3-4 trees \rightarrow red-black trees.

Red-Black Trees

110/129

Red-black trees are a representation of 2-3-4 trees using BST nodes.

- each node needs one extra value to encode link type
- but we no longer have to deal with different kinds of nodes

Link types:

- *red* links ... combine nodes to represent 3- and 4-nodes
- *black* links ... analogous to "ordinary" BST links (child links)

Advantages:

- standard BST search procedure works unmodified
- get benefits of 2-3-4 tree self-balancing (although deeper)

... Red-Black Trees

111/129

Definition of a *red-black tree*

- a BST in which each node is marked red or black
- no two red nodes appear consecutively on any path
- a red node corresponds to a 2-3-4 *sibling* of its parent
- a black node corresponds to a 2-3-4 *child* of its parent
 - if no parent (= root) → also black

Balanced red-black tree

- all paths from root to leaf have same number of black nodes


Insertion algorithm: avoids worst case $O(n)$ behaviour

Search algorithm: standard BST search

... Red-Black Trees

112/129

Representing 4-nodes in red-black trees:




Some texts colour the links rather than the nodes.

... Red-Black Trees

113/129

Representing 3-nodes in red-black trees (two possibilities):



... Red-Black Trees

114/129

Equivalent trees (one 2-3-4, one red-black):



... Red-Black Trees

115/129


Red-black tree implementation:

```
typedef enum {RED, BLACK} Colr;
typedef struct node *RBTree;
typedef struct node {
    Item data; // actual data
    Colr color; // relationship to parent
    RBTree left; // left subtree
    RBTree right; // right subtree
} node;
```

```
#define color(tree) ((tree)->color)
#define isRed(tree) ((tree) != NULL && (tree)->color == RED)
```

RED = node is part of the same 2-3-4 node as its parent (sibling)



BLACK = node is a child of the 2-3-4 node containing the parent



Exercise #14: Red-Black Trees

116/129

Show a red-black tree that corresponds to this 2-3-4 tree:



... Red-Black Trees

118/129

Search method is standard BST search:

SearchRedBlack(tree,item):

```

Input  tree, item
Output true if item found in red-black tree
        false otherwise

if tree is empty then
    return false
else if item < data(tree) then
    return SearchRedBlack(left(tree), item)
else if item > data(tree) then
    return SearchRedBlack(right(tree), item)
else
    return true // found
end if

```

Red-Black Tree Insertion

119/129

Insertion is more complex than for standard BSTs

- splitting/promoting implemented by rotateLeft/rotateRight
- several cases to consider depending on colour/direction combinations

New nodes are always **red** by default:

```

RBTree newNode(Item it) {
    RBTree new = malloc(sizeof(Node));
    assert(new != NULL);
    data(new) = it;
    colour(new) = RED;
    left(new) = right(new) = NULL;
    return new;
}

```

... Red-Black Tree Insertion

120/129

High-level description of insertion algorithm:

insertRedBlack(tree,item):

```

Input  red-black tree, item
Output tree with item inserted

```

```

tree = insertRB(tree, item)
color(tree) = BLACK // root node is always black
return tree

```

insertRB(tree,item):

```

Input  tree, item
Output tree with it inserted

```

```


if tree is empty then
    return newNode(item)
else if item = data(tree) then
    return tree
end if
if tree is a 4-node then
    split 4-node
end if
recursive insert a la BST, re-arrange links/colours after insert
return modified tree

```

... Red-Black Tree Insertion

121/129

Splitting a 4-node, in a red-black tree:



Algorithm:

```

color(left(currentTree)) = BLACK
color(right(currentTree)) = BLACK
color(currentTree) = RED

```

... Red-Black Tree Insertion

122/129

Simple recursive insert (a la BST):



Algorithm:

```
if item < data(tree) then
    left(tree) = insertRB(left(tree), item)
    re-arrange links/colours after insert
else // item larger than data in root
    right(tree) = insertRB(right(tree), item)
    re-arrange links/colours after insert
end if
```

Not affected by colour of tree node.

... Red-Black Tree Insertion

123/129

Re-arrange links/colours after insert:

Step 1 — "normalise" direction of two consecutive red nodes after insert

Algorithm:

```
if both left child and left-right grandchild of t are red then
    left-rotate left(t)
end if
```

Symmetrically,

- if both right child and right-left grandchild of t are red
⇒ right-rotate right(t)




This is in preparation for step 2 ...

... Red-Black Tree Insertion

124/129

Re-arrange links/colours after insert:

Step 2 — two consecutive red nodes = newly-created 4-node



Algorithm:

```
if both left child and left-left grandchild are red then
    t = rotateRight(t)
    color(t) = BLACK
    color(right(t)) = RED
end if
```

Symmetrically,

- if both right child and right-right grandchild are red
⇒ left rotate t, then re-colour current tree t and left(t)

... Red-Black Tree Insertion

125/129

Example of insertion, starting from empty tree:

22, 12, 8, 15, 11, 19, 43, 44, 45, 42, 41, 40, 39



Red-black Tree Performance

126/129

Cost analysis for red-black trees:

- tree is well-balanced; worst case search is $O(\log_2 n)$
- insertion affects nodes down one path; #rotations+recolorings is $O(h)$
(where h is the height of the tree)

Only disadvantage is complexity of insertion/deletion code.

Note: red-black trees were popularised by Sedgewick.

Application of BSTs: Sets

127/129

Trees provide efficient search.

Sets require efficient search

- to find where to insert/delete
- to test for set membership

Logical to implement a Set ADT via `BSTree`

Assuming we have `BSTree` implementation

- which precludes duplicate key values
- which implements insertion, search, deletion

then `Set` implementation is

- `addToSet(Set, Item) ≡ TreeInsert(Tree, Item)`
- `removeFromSet(Set, Item) ≡ TreeDelete(Tree, Item.Key)`
- `elementOfSet(Set, Item) ≡ TreeSearch(Tree, Item.Key)`

... Application of BSTs: Sets


128/129

Concrete representation:


```
#include "BSTree.h"
```

```
typedef struct SetRep {
    int    nelems;
    Tree  root;
} SetRep;
```


```
typedef SetRep *Set;
```



After `SetInsert(s,4)`:



After `SetDelete(s,2)`:



Summary

129/129

- Binary search tree (BST) data structure
- Tree traversal

- Tree operations
 - insertion, join, deletion, rotation
 - tree partition, rebalancing
- Self-adjusting trees
 - Splay trees
 - AVL trees
 - 2-3-4 trees
 - Red-black trees

- Suggested reading (Sedgewick):
 - BSTs ... Ch. 12.5-12.6
 - rotation, partition, deletion, join ... Ch. 12.8-12.9
 - self-adjusting trees ... Ch. 13.1-13.4

Produced: 27 Jan 2025