

COMP9313: Big Data Management



Lecturer: Xubo Wang

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 4.1: Spark I

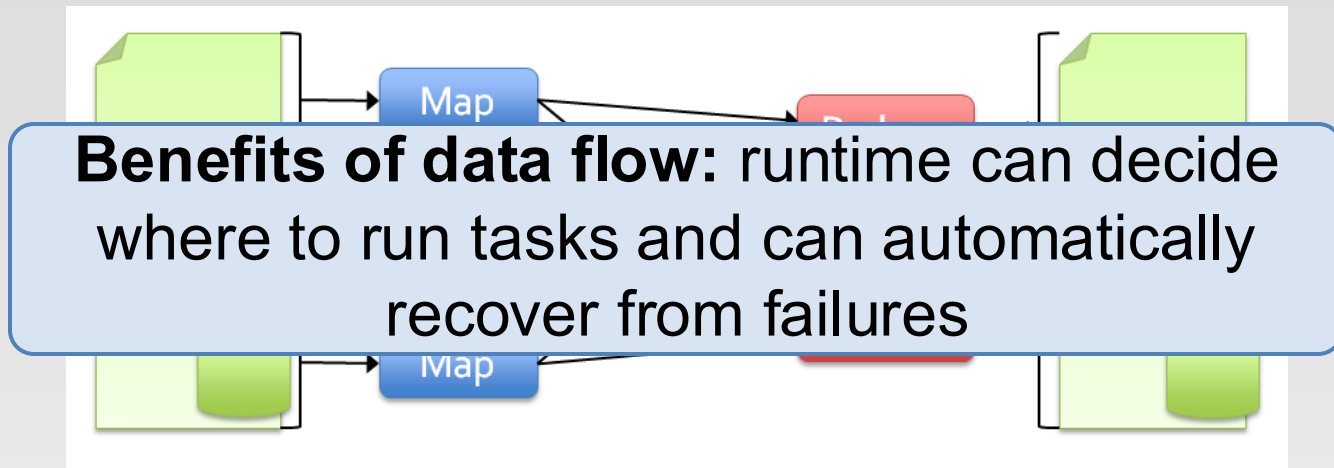


Part 1: Spark Introduction

Limitations of MapReduce

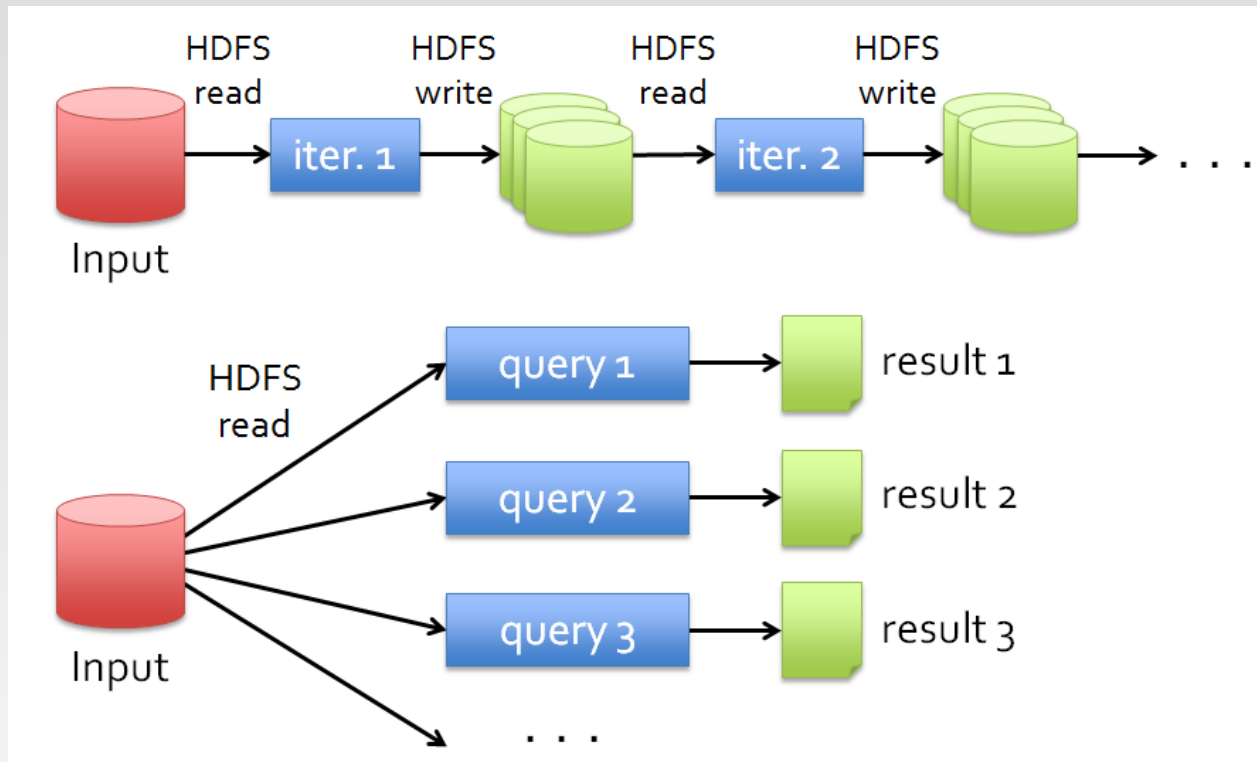
- ❖ MapReduce greatly simplified big data analysis on large, unreliable clusters. It is great at one-pass computation.
- ❖ But as soon as it got popular, users wanted more:
 - More **complex**, multi-pass analytics (e.g. ML, graph)
 - More **interactive** ad-hoc queries
 - More **real-time** stream processing
- ❖ All 3 need faster **data sharing** across parallel jobs
 - One reaction: specialized models for some of these apps, e.g.,
 - ▶ Pregel (graph processing)
 - ▶ Storm (stream processing)

Limitations of MapReduce



- ❖ As a general programming model:
 - It is more suitable for one-pass computation on a large dataset
 - Hard to compose and nest multiple operations
 - No means of expressing iterative operations
- ❖ As implemented in Hadoop
 - All datasets are read from disk, then stored back on to disk
 - All data is (usually) triple-replicated for reliability
 - Not easy to write MapReduce programs using Java

Data Sharing in MapReduce




Slow due to replication, serialization, and disk IO

- ❖ Complex apps, streaming, and interactive queries all need one thing that MapReduce lacks:


Efficient primitives for **data sharing**

Data Sharing in MapReduce

- ❖ Iterative jobs involve a lot of disk I/O for each repetition



- ❖ Interactive queries and online processing involves lots of disk I/O



Hardware for Big Data



Lots of hard drives



Lots of CPUs




And lots of memory!

Goals of Spark

- ❖ Keep more data in-memory to improve the performance!
- ❖ Extend the MapReduce model to better support two common classes of analytics apps:
 - Iterative algorithms (machine learning, graphs)
 - Interactive data mining
- ❖ Enhance programmability:
 - Integrate into Scala programming language
 - Allow interactive use from Scala interpreter


Data Sharing in Spark Using RDD



10-100 × faster than network and disk

What is Spark

- ❖ One popular answer to “What’s beyond MapReduce?”
- ❖ Open-source engine for large-scale distributed data processing
 - Supports generalized dataflows
 - Written in Scala, with bindings in Java, Python, and R
- ❖ Brief history:
 - Developed at UC Berkeley AMPLab in 2009
 - Open-sourced in 2010
 - Became top-level Apache project in February 2014
 - Commercial support provided by DataBricks



What is Spark

- ❖ Fast and expressive cluster computing system interoperable with Apache Hadoop
- ❖ Improves efficiency through:
 - **In-memory** computing primitives
 - General computation graphs

→ Up to 100 × faster
(10 × on disk)
- ❖ Improves usability through:
 - Rich APIs in Scala, Java, Python
 - Interactive shell

→ Often 5 × less code

What is Spark

❖ Spark is not

- a modified version of Hadoop
- dependent on Hadoop because it has its own cluster management
- Spark uses Hadoop for storage purpose only

❖ Spark's design philosophy centers around four key characteristics:

- Speed
- Ease of use
- Modularity
- Extensibility

Speed

- ❖ Its internal implementation benefits immensely from the performance improvement of CPUs and memory.
 - The framework is optimized to take advantage of memory, multiple cores, and the underlying Unix-based operating system
- ❖ Spark builds its query computations as a Directed Acyclic Graph
 - Tasks can execute in parallel across workers on the cluster
- ❖ It has a physical execution engine which generates compact code for execution

Ease of Use

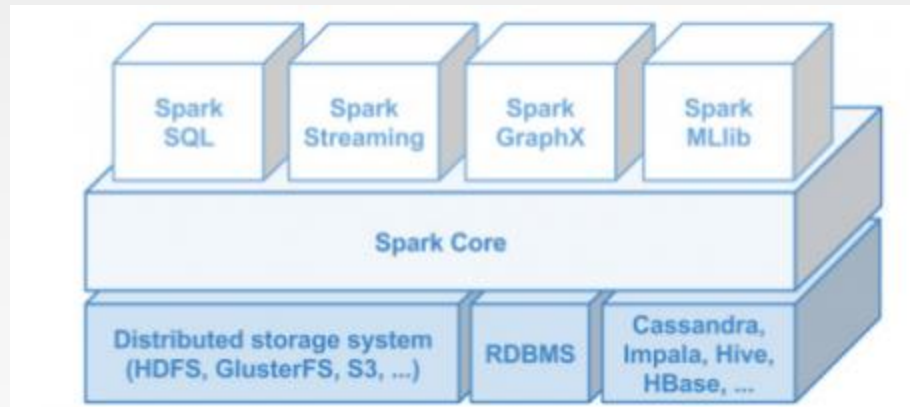
- ❖ Spark achieves simplicity by providing a fundamental abstraction of a simple logical data structure called a Resilient Distributed Dataset (RDD)
- ❖ Since Spark 2.x, DataFrames and Datasets APIs have been developed upon RDD
- ❖ By providing a set of ***transformations*** and ***actions*** as operations, Spark offers a simple programming model that you can use to build big data applications in familiar languages.

Modularity

- ❖ Spark operations can be applied across many types of workloads and expressed in any of the supported programming languages: Scala, Java, Python, SQL, and R.
- ❖ Spark offers unified libraries with well-documented APIs that include the following modules as core components: Spark SQL, Spark Structured Streaming, Spark MLlib, and GraphX, combining all the workloads running under one engine.
- ❖ You can write a single Spark application that can do it all—no need for distinct engines for disparate workloads, no need to learn separate APIs.

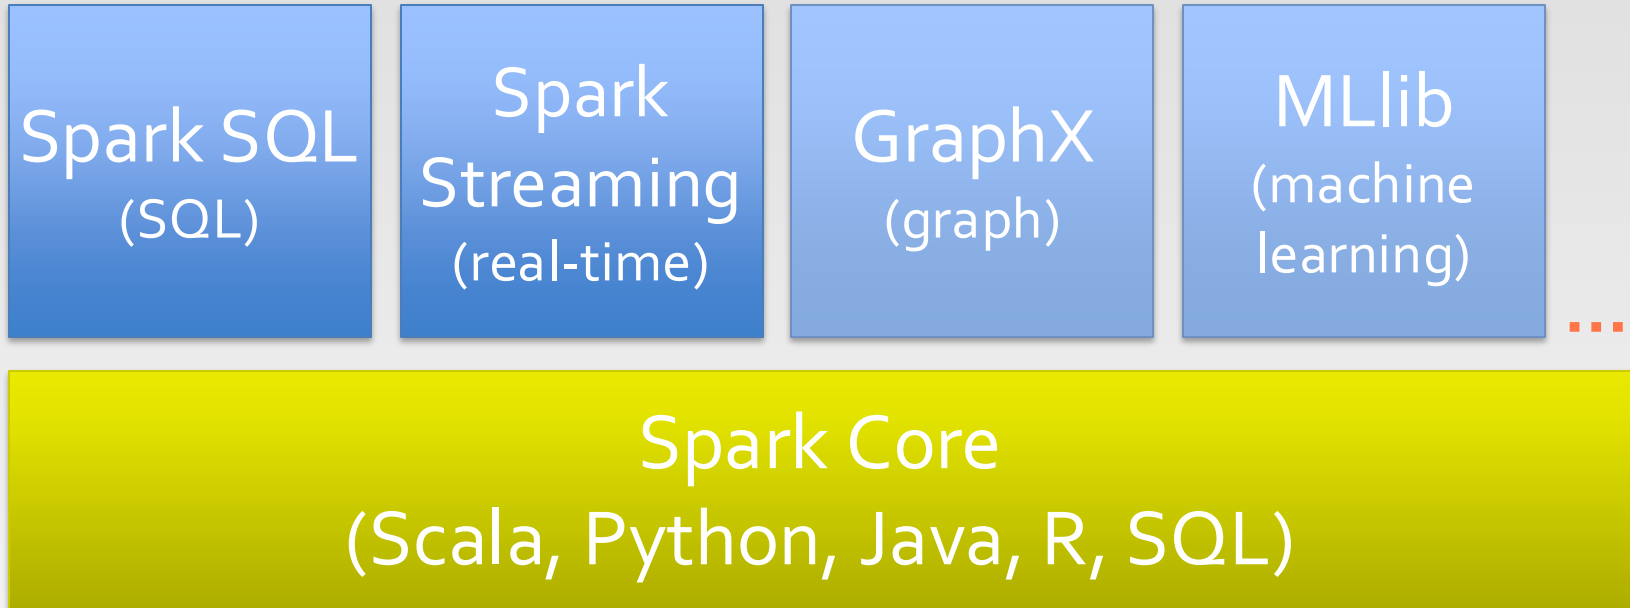
Extensibility

- ❖ Spark focuses on its fast, parallel computation engine rather than on storage.
 - You can use Spark to read data stored in myriad sources—local file systems, Apache Hadoop, Apache Cassandra, Apache HBase, MongoDB, Apache Hive, RDBMSs, and more—and process it all in memory.
- ❖ Spark's DataFrameReaders and DataFrameWriters can also be extended to read data from other sources, such as Apache Kafka, Kinesis, Azure Storage, and Amazon S3



What is Spark


- ❖ Spark is the basis of a wide set of projects in the Berkeley Data Analytics Stack (BDAS)



- Spark SQL (SQL on Spark)
- Spark Streaming (stream processing)
- GraphX (graph processing)
- MLlib (machine learning library)

Spark's Ecosystem of Connectors

- ❖ The community of Spark developers maintains a list of third-party Spark packages as part of the growing ecosystem

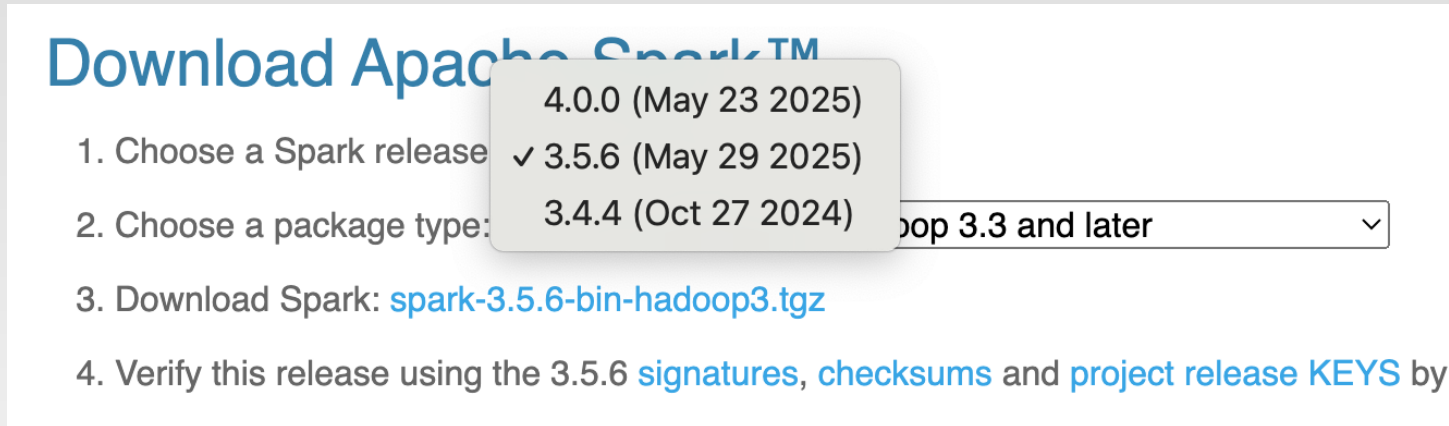


Spark Ideas

- ❖ Expressive computing system, not limited to map-reduce model
- ❖ Facilitate system memory
 - avoid saving intermediate results to disk
 - cache data for repetitive queries (e.g. for machine learning)
- ❖ Layer an in-memory system on top of Hadoop.
- ❖ Achieve fault-tolerance by re-execution instead of replication

Download and Configure Spark

- ❖ Current version: 3.5.6. <https://spark.apache.org/downloads.html>
 - You also need to install Java first



The screenshot shows the 'Download Apache Spark™' page. It features a list of steps for downloading Spark. A dropdown menu is open for 'Choose a Spark release', showing three options: 4.0.0 (May 23 2025), 3.5.6 (May 29 2025) which is selected with a checkmark, and 3.4.4 (Oct 27 2024). To the right of the dropdown is a text input field containing 'hadoop 3.3 and later' with a dropdown arrow. Below the dropdown, step 3 indicates the download file 'spark-3.5.6-bin-hadoop3.tgz'. Step 4 instructs to verify the release using signatures, checksums, and project release keys.

Download Apache Spark™

1. Choose a Spark release: 4.0.0 (May 23 2025) ✓ 3.5.6 (May 29 2025) 3.4.4 (Oct 27 2024)
2. Choose a package type: hadoop 3.3 and later
3. Download Spark: [spark-3.5.6-bin-hadoop3.tgz](#)
4. Verify this release using the 3.5.6 [signatures](#), [checksums](#) and [project release KEYS](#) by

- ❖ After downloading the package, unpack it and then configure the path variable in file ~/.bashrc

```
export SPARK_HOME=/home/comp9313/spark
export PATH=$SPARK_HOME/bin:$SPARK_HOME/sbin:$PATH
```

Spark Shell

- ❖ Spark comes with four widely used interpreters that act like interactive “shells” and enable ad hoc data analysis: pyspark, spark-shell, sparksql, and sparkR

```
[192-168-1-100:9313 xubowang$ pyspark
Python 3.11.4 (v3.11.4:d2340ef257, Jun  6 2023, 19:15:51) [Clang 13.0.0 (clang-1300.0.29.30)] on
darwin
Type "help", "copyright", "credits" or "license" for more information.
24/06/17 15:47:21 WARN Utils: Your hostname, 192-168-1-100.tpgi.com.au resolves to a loopback add
ress: 127.0.0.1; using 192.168.1.100 instead (on interface en0)
24/06/17 15:47:21 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
24/06/17 15:47:22 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform..
. using builtin-java classes where applicable
Welcome to

  ____      _
 / ___|    / \
| |  | |  / _ \
| |  | | / ___ \
| |  | | \___/ \
| |  | |
| |  | |
|_|  |_|

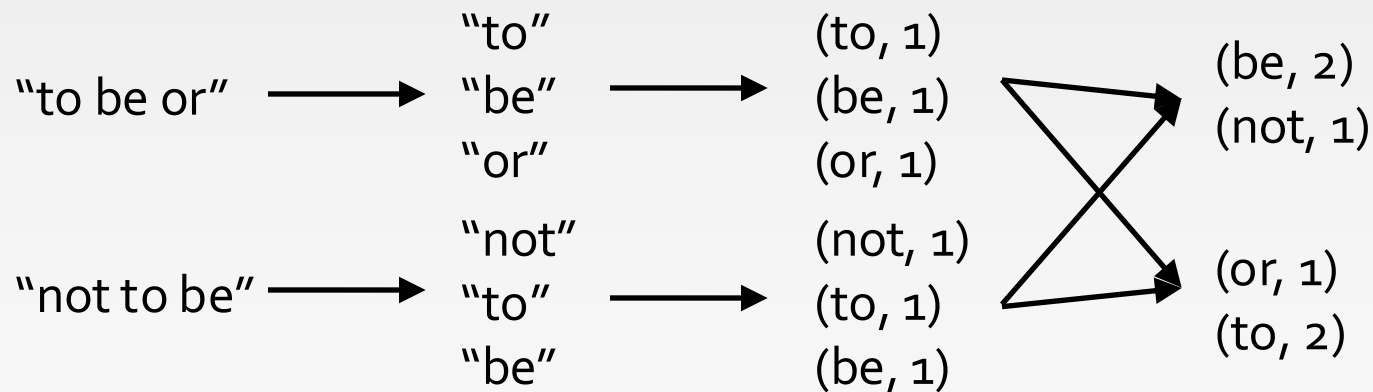
 version 3.5.1

Using Python version 3.11.4 (v3.11.4:d2340ef257, Jun  6 2023 19:15:51)
Spark context Web UI available at http://192.168.1.100:4040
Spark context available as 'sc' (master = local[*], app id = local-1718603243664).
SparkSession available as 'spark'.
[>>>
[>>>
```

Word Count in Spark

```
textfile = sc.textFile("hdfs://...", 4)

words = textfile.flatMap(Lambda line: line.split(" "))
pairs = words.map(Lambda word: (word, 1))
count = pairs.reduceByKey(Lambda a, b: a + b)
count.collect()
```



Python Supports Functional Programming

- ❖ First-Class Functions. Functions are treated like objects:
 - passing functions as arguments to other functions
 - returning functions as the values from other functions
 - assigning functions to variables or storing them in data structures

```
>>> plusOne = lambda x: x+1
>>> type(plusOne)
<class 'function'>
>>> plusOne(5)
6
```

```
>>> def f(plusOne):
...     return plusOne(5)
...
>>> f(plusOne)
6
```

```
>>> def f(x):
...     return lambda y:y+x
...
>>> f(5)(2)
7
```


Closures

- ❖ Closures: a function whose return value depends on the value of one or more variables declared outside this function.

// plusFoo can reference any **values/variables** in scope

foo = 1

plusFoo = lambda x: x+foo

plusFoo(5) → 6

// Changing foo changes the return value of plusFoo

foo = 5

plusFoo(5) → 10

Higher Order Functions

❖ Higher Order Functions

- A function that does at least one of the following:
 - ▶ takes one or more functions as arguments
 - ▶ returns a function as its result

```
>>> def f(plusOne):  
...     return plusOne(5)  
...  
>>> f(plusOne)  
6
```

```
>>> def f(x):  
...     return lambda y:y+x  
...  
>>> f(5)(2)  
7
```

More Examples on Higher Order Functions

```
basefunc = lambda x: (lambda y: x + y)
```

```
// interpreted by:
```

```
basefunc(x):  
    sumfunc(y):  
        return x+y  
    return sumfunc
```

```
closure1 = basefunc(1)           closure1(5) = ?  
                                     6
```

```
closure2 = basefunc(4)           closure2(5) = ?  
                                     9
```

- ❖ basefunc returns a function, and closure1 and closure2 are of function type.
- ❖ While closure1 and closure2 refer to the same function basefunc, the associated environments differ, and the results are different

Part 2: RDD Introduction

RDD: Resilient Distributed Datasets

- ❖ Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. Matei Zaharia, et al. NSDI'12
 - RDD is a **distributed** memory abstraction that lets programmers perform **in-memory** computations on large clusters in a **fault-tolerant** manner.
- ❖ **Resilient**
 - Fault-tolerant, is able to recompute missing or damaged partitions due to node failures.
- ❖ **Distributed**
 - Data residing on multiple nodes in a cluster.
- ❖ **Dataset**
 - A collection of partitioned elements, e.g. tuples or other objects (that represent records of the data you work with).
- ❖ RDD is the primary data abstraction in Apache Spark and the core of Spark. It enables operations on collection of elements in parallel.

RDD: Resilient Distributed Datasets

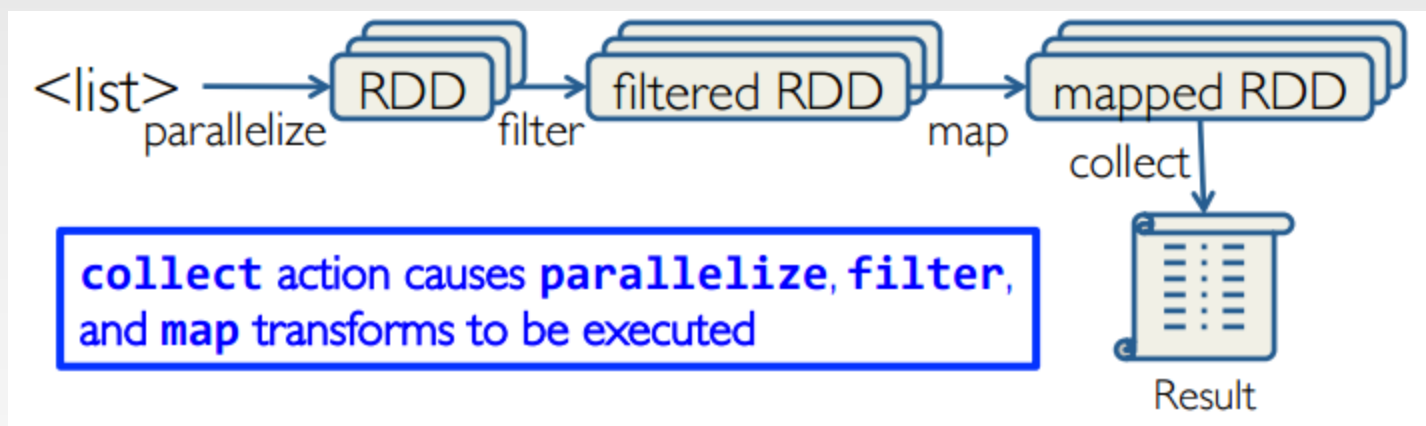
- ❖ *Resilient Distributed Datasets (RDDs)*
 - Distributed collections of objects that can be cached in memory across cluster
 - Manipulated through parallel operators
 - Automatically recomputed on failure based on lineage
- ❖ RDDs can express many parallel algorithms, and capture many current programming models
 - Data flow models: MapReduce, SQL, ...
 - Specialized models for iterative apps: Pregel, ...

RDD Traits

- ❖ **In-Memory**, i.e. data inside RDD is stored in memory as much (size) and long (time) as possible.
- ❖ **Immutable** or **Read-Only**, i.e. it does not change once created and can only be transformed using transformations to new RDDs.
- ❖ **Lazy evaluated**, i.e. the data inside RDD is not available or transformed until an action is executed that triggers the execution.
- ❖ **Cacheable**, i.e. you can hold all the data in a persistent "storage" like memory (default and the most preferred) or disk (the least preferred due to access speed).
- ❖ **Parallel**, i.e. process data in parallel.
- ❖ **Typed**, i.e. values in a RDD have types, e.g. RDD[Long] or RDD[(Int, String)].
- ❖ **Partitioned**, i.e. the data inside a RDD is partitioned (split into partitions) and then distributed across nodes in a cluster (one partition per JVM that may or may not correspond to a single node).

Working with RDDs

- ❖ Create an RDD from a data source
 - by parallelizing existing collections (lists or arrays)
 - by transforming an existing RDDs
 - from files in HDFS or any other storage system
- ❖ Apply transformations to an RDD: e.g., map, filter
- ❖ Apply actions to an RDD: e.g., collect, count



- ❖ Users can control two other aspects:
 - Persistence
 - Partitioning

Creating RDDs

- ❖ From HDFS, text files, Amazon S3, Apache HBase, SequenceFiles, any other Hadoop InputFormat
- ❖ Creating an RDD from a File
 - `inputfile = sc.textFile("...", 4)`
 - ▶ RDD distributed in 4 partitions
 - ▶ Elements are lines of input
 - ▶ Lazy evaluation means no execution happens now

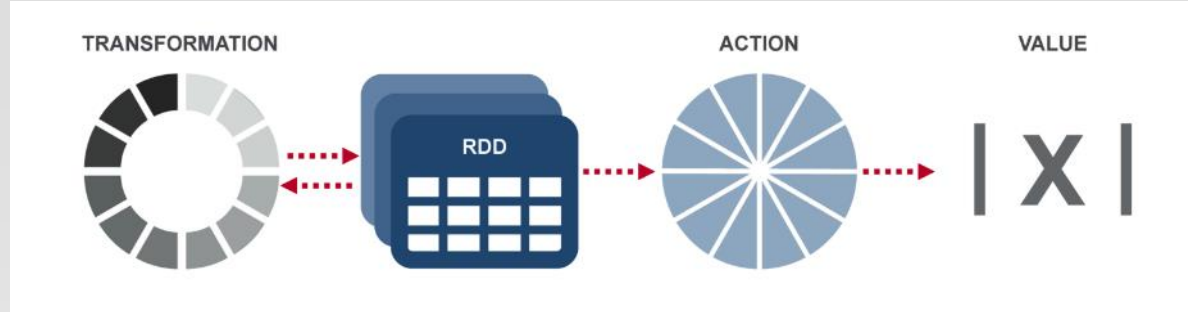
```
>>> inputfile = sc.textFile("file:///home/comp9313/pg100.txt")
>>> inputfile
file:///home/comp9313/pg100.txt MapPartitionsRDD[1] at textFile at NativeMethodAccessImpl.java:0
```

- ❖ Turn a collection into an RDD
 - `sc.parallelize([1, 2, 3])`, creating from a Python list

Repartition and Coalesce

- ❖ Sometimes we may need to repartition the RDD, PySpark provides two ways to repartition:
 - `repartition()`: shuffles data from all nodes, also called full shuffle
 - `coalesce()`: shuffle data from minimum nodes. For example, if you have data in 4 partitions, doing `coalesce(2)` moves data from just 2 nodes.
 - Both of the functions take the number of partitions to repartition rdd.
 - Note that `repartition()` method is a very expensive operation as it shuffles data from all nodes in a cluster.
 - `repartition()` is used to increase or decrease the RDD partitions whereas `coalesce()` is used to **only** decrease the number of partitions in an efficient way.

RDD Operations



- ❖ **Transformation:** returns a new RDD.
 - Nothing gets evaluated when you call a Transformation function, it just takes an RDD and return a new RDD.
 - Transformation functions include *map*, *filter*, *flatMap*, *groupByKey*, *reduceByKey*, *aggregateByKey*, *join*, etc.
- ❖ **Action:** evaluates and returns a new value.
 - When an Action function is called on a RDD object, all the data processing queries are computed at that time and the result value is returned.
 - Action operations include *reduce*, *collect*, *count*, *first*, *take*, *countByKey*, *foreach*, *saveAsTextFile*, etc.
- ❖ <https://spark.apache.org/docs/3.5.6/api/python/reference/api/pyspark.RDD.html#pyspark.RDD>

Spark Transformations

- ❖ Create new datasets from an existing one
- ❖ Use lazy evaluation: results not computed right away – instead Spark remembers set of transformations applied to base dataset
 - Spark optimizes the required calculations
 - Spark recovers from failures
- ❖ Some transformation functions

Transformation	Description
<code>map(func)</code>	return a new distributed dataset formed by passing each element of the source through a function <i>func</i>
<code>filter(func)</code>	return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true
<code>distinct([numTasks])</code>	return a new dataset that contains the distinct elements of the source dataset
<code>flatMap(func)</code>	similar to map, but each input item can be mapped to 0 or more output items (so <i>func</i> should return a Seq rather than a single item)

Spark Actions

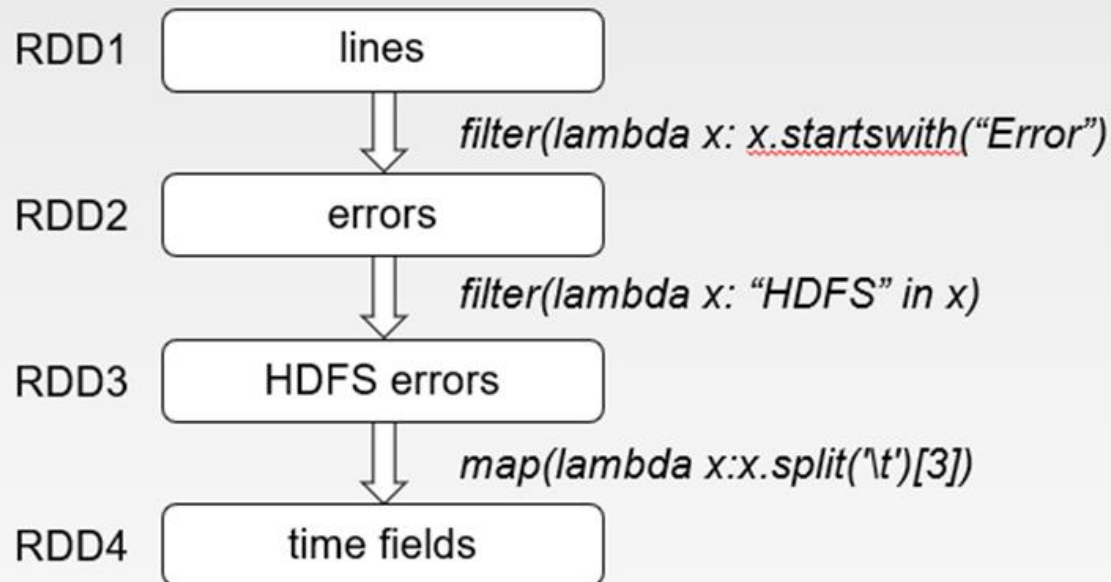
- ❖ Cause Spark to execute recipe to transform source
- ❖ Mechanism for getting results out of Spark
- ❖ Some action functions

Action	Description
<code>reduce(func)</code>	aggregate dataset's elements using function <i>func</i> . <i>func</i> takes two arguments and returns one, and is commutative and associative so that it can be computed correctly in parallel
<code>take(n)</code>	return an array with the first <i>n</i> elements
<code>collect()</code>	return all the elements as an array WARNING: make sure will fit in driver program
<code>takeOrdered(n, key=func)</code>	return <i>n</i> elements ordered in ascending order or as specified by the optional key function

Example

- ❖ Web service is experiencing errors and an operator wants to search terabytes of logs in the Hadoop file system to find the cause.

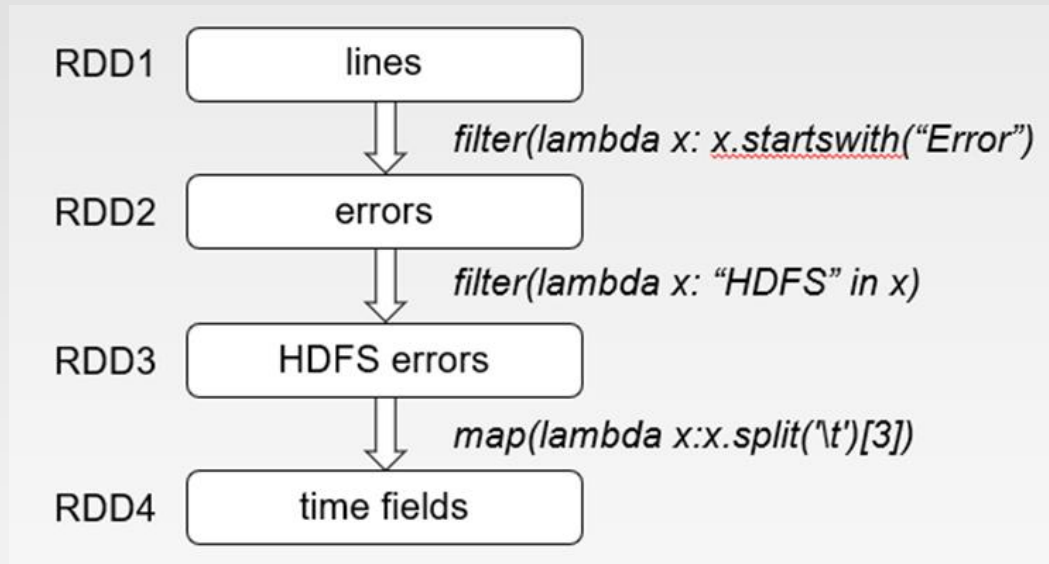
```
lines = sc.textFile("hdfs://...") //base RDD, obtained from a file on HDFS
errors = lines.filter(lambda x: x.startswith("Error")) //get messages that start
errors.persist() //persist the data in memory
errors.count()
errors.filter(lambda x: "HDFS" in x).map(lambda x:x.split('\t')[3]).collect()
```



- **Line1:** RDD backed by an HDFS file (base RDD lines not loaded in memory)
- **Line3:** Asks for errors to persist in memory (errors are in RAM)


Lineage Graph

- ❖ RDDs keep track of lineage
- ❖ RDD has enough information about how it was derived from to compute its partitions from data in stable storage.



- ❖ Example:
 - If a partition of errors is lost, Spark rebuilds it by applying a filter on only the corresponding partition of lines.
 - Partitions can be recomputed in parallel on different nodes, without having to roll back the whole program.

Deconstructed



//base RDD

lines = sc.textFile("hdfs://...")

//Transformed RDD

errors = lines.filter(lambda x: x.startswith("Error"))

errors.persist()

errors.count()

errors.filter(lambda x: "HDFS" in x).

map(lambda x: x.split('\t')[3]).


collect()

Deconstructed



//base RDD


lines = sc.textFile("hdfs://...")



//Transformed RDD

errors = lines.filter(lambda x: x.startswith("Error"))

errors.persist()



count() causes Spark to: 1) read data; 2) sum within partitions; 3) combine sums in driver

Put transform and action together:

errors.filter(lambda x: "HDFS" in x).map(lambda x: x.split("\t")[3]).collect()

RDD Persistence: Cache/Persist

- ❖ One of the most important capabilities in Spark is *persisting* (or *caching*) a dataset in memory across operations.
- ❖ When you persist an RDD, each node stores any partitions of it. You can reuse it in other actions on that dataset
- ❖ Each persisted RDD can be stored using a different *storage level*, e.g.
 - MEMORY_ONLY:
 - ▶ Store RDD as deserialized Java objects in the JVM.
 - ▶ If the RDD does not fit in memory, some partitions will not be cached and will be recomputed when they're needed.
 - ▶ This is the default level.
 - MEMORY_AND_DISK:
 - ▶ If the RDD does not fit in memory, store the partitions that don't fit on disk, and read them from there when they're needed.
- ❖ `cache()` = `persist(StorageLevel.MEMORY_ONLY)`

Why Persisting RDD?

```
lines = sc.textFile("hdfs://...")
```

```
errors = lines.filter(lambda x: x.startswith("Error"))
```

```
errors.persist()
```

```
errors.count()
```

- ❖ If you do `errors.count()` again, the file will be loaded again and computed again.
- ❖ `Persist` will tell Spark to cache the data in memory, to reduce the data loading cost for further actions on the same data
- ❖ `errors.persist()` will do nothing. It is a lazy operation. But now the RDD says "read this file and then cache the contents". The action will trigger computation and data caching.

References

- ❖ <http://spark.apache.org/docs/latest/index.html>
- ❖ [Learning Spark](#). 1st and 2nd Edition

End of Chapter 4.1