COMP9444: Neural Networks and Deep Learning

Week 1b. Neuroanatomy

Alan Blair

School of Computer Science and Engineering

May 26, 2025

What is a Neural Network?

- ➤ massively parallel distributed processor made up of simple processing units
- ➤ knowledge acquired from environment through a learning process
- ➤ knowledge stored in the form of synaptic weights

Why Neural Networks?

- ➤ biologically inspired
- ➤ good learning properties
- ➤ continuous, nonlinear
- ➤ well adapted to certain tasks
- ➤ fault tolerant
- ➤ graceful degradation

UNSW

Sub-Symbolic Processing

Dot Illusion

Theories about Intelligence

- 380BC Plato (Rationalism - innateness)
- 330BC Aristotle (Empricism - experience)
- 1641 Descartes (mind-body Dualism)
- 1781 Kant (Critique of Pure Reason)
- 1899 Sigmund Freud (Psychology)
- 1953 B.F. Skinner (Behaviourism)

UNSW

Artificial Intelligence Origins

- ➤ 1642 Blaise Pascal (mechanical adding machine)
- ➤ 1694 Gottfried Leibniz (mechanical calculator)
- ➤ 1769 Wolfgang von Kempelen (Mechanical Turk)
- ➤ 1837 Charles Babbage & Ada Lovelace (Difference Engine)
- ➤ 1848 George Boole (the Calculus of Logic)
- ➤ 1879 Gottlob Frege (Predicate Logic)
- ➤ 1950 Turing Test
- ➤ 1956 Dartmouth conference

Neural Network Origins

- 1943 McCulloch & Pitts (neuron models)
- 1948 Norbert Wiener (Cybernetics)
- 1948 Alan Turing (B-Type Networks)
- 1955 Oliver Selfridge (Pattern Recognition)
- 1962 Hubel and Wiesel (visual cortex)
- 1962 Frank Rosenblatt (Perceptron)

Serial Symbolic AI

- ➤ 1956 Newell & Simon (Logic Theorist)
- ➤ 1959 John McCarthy (Lisp)
- ➤ 1959 Arther Samuel (Checkers)
- ➤ 1965 Joseph Weizenbaum (ELIZA)
- ➤ 1967 Edward Feigenbaum (Dendral)

UNSW

Neural Network "Dark Ages"

- ➤ 1969 Minsky & Papert published "Perceptrons", emphasizing the limitations of neural models, and lobbied agencies to cease funding neural network research.
- ➤ from 1969 to 1985, very little work in neural networks or machine learning.
- ➤ a few exceptions, e.g. Stephen Grossberg, Teuvo Kohonen, Paul Werbos.

Knowledge-Based Systems

- ➤ 1970's and early 1980's, AI research focused on symbolic processing, Expert Systems
- ➤ Some commercial success, but ran into difficulties:
  - → combinatorial explosion in search spaces
  - → difficulty of formalising everyday knowledge as well as expert knowledge

UNSW

Neural Network Renaissance

- → 1986 Rumelhart, Hinton & Williams (multi-layer, backprop)
- → 1989 Dean Pomerleau (ALVINN)
- → late 1980's renewed enthusiasm, hype
- → 1990's more principled approaches
- → 2000's SVM, Bayesian models became more popular
- → 2010's deep learning networks, GPU's
- → 2020's transformers, stable diffusion

Applications of Deep Learning

- ➤ Image Processing
  - → Classification
  - → Segmentation
  - → Image Generation

- ➤ Language Processing
  - → Translation
  - → Sentiment Analysis
  - → Text Generation
  - → Chatbots

- ➤ Game Playing
  - → Deep Q-Learning
  - → AlphaGo
  - → AlphaStar

- ➤ Combining Images and Language
  - → Automatic Captioning
  - → Generating Images from Text
  - → Video Generation
  - → Multimodal AI

History of Deep Learning

Two perspectives on the history of Deep Learning

Viewpoint 1: Focusing on recent work (after 2012)
`https://www.cs.toronto.edu/~hinton/absps/NatureDeepReview.pdf`

Viewpoint 2: Focusing on earlier work (before 2012)
`http://people.idsia.ch/~juergen/deep-learning-overview.html`

Neuroanatomy

- Central Nervous System
    - Brain
    - Spinal cord

- Peripheral Nervous System
    - Somatic nervous system
    - Autonomic nervous system
    - Enteric nervous system

Brain Regions

Cerebral hemisphere — Midbrain — Thalamus — Hypothalamus — Pituitary — Pons — Medulla — Cerebellum — Spinal cord

Cerebral Cortex

Frontal lobe
Parietal lobe
Occipital lobe
Left hemisphere
Cerebrum
Cerebellum
Temporal lobe
Spinal cord
Brain stem

UNSW

Cerebral Cortex

- "cortex" from Latin word for "bark" (of tree)
- cortex is a sheet of tissue making up outer layers of brain, 2-6cm thick
- right and left sides connected by corpus callosum
- functions: thought, voluntary movement, language, reasoning, perception

Brain Stem

- general term for area of brain between the thalamus and spinal cord
- includes medulla, pons, tectum, reticular formation and tegmentum
- functions: breathing, heart rate, blood pressure, and others

Cerebellum

- from Latin word for "little brain"
- functions: movement, balance, posture

UNSW

Midbrain

- ➤ functions: vision, audition, eye movement, body movement

Thalamus

- ➤ receives sensory information and relays it to the cerebral cortex
- ➤ also relays information from the cerebral cortex to other areas of the brain, and the spinal cord
- ➤ functions: sensory integration, motor integration

Hypothalamus

- ➤ composed of several different areas at the base of the brain
- ➤ the size of a pea (about 1/300 of the total brain weight)
- ➤ functions: body temperature, emotions, hunger, thirst, circadian rhythms

Limbic System

Limbic System

- group of structures including amygdala, hippocampus, mammillary bodies and cingulate gyrus
- important for controlling the emotional response to a given situation
- hippocampus also important for memory
- functions: emotional behaviour

Brain Functions

Neurons as Body Cells

- ➤ The body is made up of billions of cells. Cells of the nervous system, called *neurons*, are specialized to carry "messages" through an electrochemical process.
- ➤ The human brain has about 100 billion neurons, and a similar number of support cells called "glia".
- ➤ Neurons are similar to other cells in the body in some ways, such as:
  - → neurons are surrounded by a cell membrane
  - → neurons have a nucleus that contains genes (DNA)
  - → neurons carry out basic cellular processes like protein synthesis and energy production

Structure of a Typical Neuron

Neurons versus Body Cells

➤ Neurons have specialized extensions called *dendrites* and *axons* Dendrites bring information to the cell body, while axons take information away from the cell body.

➤ The axon of one neuron can connect to the dendrite of another neuron through an electrochemical junction called a *synapse*.

➤ Most neurons have only one axon, but the number of dendrites can vary widely:
  → Unipolar and Bipolar neurons have only one dendrite
 → Purkinje neurons can have up to 100,000 dendrites

Variety of Neuron Types

MOTOR NEURON FROM SPINAL CORD — Dendrite, Axon

MITRAL CELL FROM OLFACTORY BULB — Dendrite, Cell body, Axon

PYRAMIDAL CELL FROM CORTEX — Dendrite, Cell body, Axon

PURKINJE CELL — Axon, Front, Side

Axons and Dendrites

- ➤ Dendrites are typically less than a millimetre in length
- ➤ Axons can vary in length from less than a millimetre to more than a metre (motor neurons)
- ➤ Long axons are sometimes surrounded by a myelinated sheath, which prevents the electrical signal from dispersing, and allows it to travel faster (up to 100 m/s).

Synapse

Voltage-gated Ca++ channels

Synaptic vesicle

Neurotransmitters

Neurotransmitter re-uptake pump

Post-synaptic density

Neurotransmitter receptors

Axon Terminal

Synaptic Cleft

Dendritic Spine

Synapses and Ion Channels

- ➤ electrical pulse reaches the endbulb and causes the release of neurotransmitter molecules from little packets (vesicles) through the synaptic membrane
- ➤ transmitter then diffuses through the synaptic cleft to the other side
- ➤ when the neurotransmitter reaches the post-synaptic membrane, it causes a change in polarisation of the membrane
- ➤ the change in potential can be *excitatiory* (moving the potential towards the threshold) or *inhibitory* (moving it away from the threshold)

Ion Channel

The Big Picture

- ➤ human brain has 100 billion neurons with an average of $10,000$ synapses each
- ➤ latency is about 3-6 milliseconds
- ➤ therefore, at most a few hundred "steps" in any mental computation, but massively parallel

Hubel and Weisel – Visual Cortex

- → cells in the visual cortex respond to lines at different angles
- → cells in V2 respond to more sophisticated visual features
- → Convolutional Neural Networks are inspired by this neuroanatomy
- → CNN's can now be simulated with massive parallelism, using GPU's

Convolutional Networks

Suppose we want to classify an image as a bird, sunset, dog, cat, etc.
If we can identify features such as feather, eye, or beak which provide useful information in one part of the image, then those features are likely to also be relevant in another part of the image.
We can exploit this regularity by using a convolution layer which applies the same weights to different parts of the image.
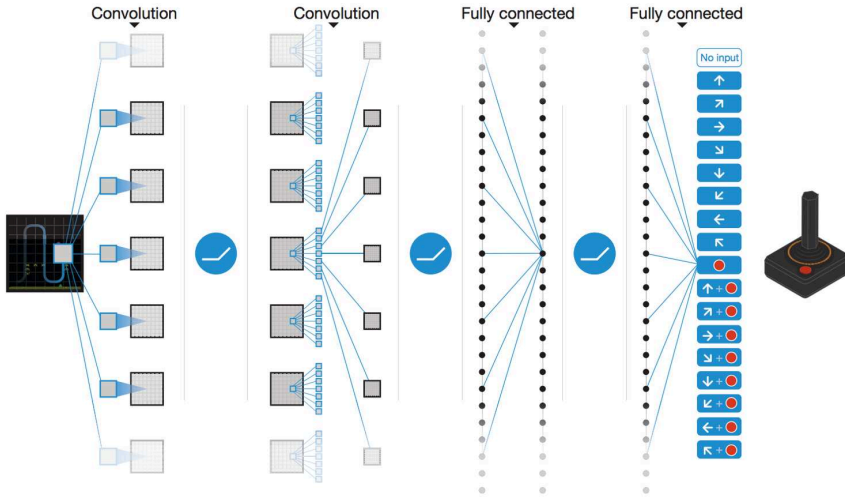
# Convolutional Filters

First Layer   **Second Layer**   **Third Layer**

Deep Q-Learning

Starcraft

Autoencoder Networks

"bottleneck" hidden layer

input layer

output layer
(reconstruction of input layer)

all layers are fully connected but not
drawn

➤ output is trained to reproduce the input as closely as possible
➤ activations normally pass through a bottleneck, so the network is forced to compress the data in some way

Generative Adversarial Networks

UNSW

Stable Diffusion

Recurrent Neural Networks

➤ can "unroll" a recurrent architecture into an equivalent feedforward architecture, with shared weights

➤ useful for processing language or other temporal sequences

Transformer